

FIT1043 Introduction to Data Science

Assignment 1

Lim Yu-Shan

32685467

30th March 2022

Introduction

This is an analysis of life expectancy, population and gross domestic product (GDP) in Southeast Asian countries. The data for life expectancy covers the years from 2000 to 2015 whereas for population and GDP, data from 2019 will be used.

All data used is sourced from the following three dataset:

- `LifeExpectancyData-v2.csv` , containing information related to life expectancy and health factors for 193 countries (source: <https://www.kaggle.com/kumarajarshi/life-expectancy-who>)
- `2019-GDP.csv` , containing information on the gross domestic product (GDP) for almost all countries in the world for the year 2019 (source: <https://datacatalog.worldbank.org/dataset/gdp-ranking>)
- `2020-Population.csv` , containing information on country and region populations from 1950 to 2020 (source: <https://population.un.org/wpp/Download/Standard/CSV/>)

Importing libraries

The very first step is to import **NumPy** and **pandas**. **NumPy** is a Python library that allows us to create special objects called arrays as well as matrices, and perform mathematical operations on these objects in addition to other useful mathematical features. **pandas**, on the other hand, is an extremely useful library for Python that provides tools for data analysis. By using **pandas**, we are able to directly convert CSV files into clear and concise data structures called *DataFrames*.

For data visualisation, we import the libraries **Matplotlib** and **Seaborn**. **Matplotlib** provides useful tools to visualise data as various graphs and charts. **Seaborn** is another data visualisation library built on top of **Matplotlib**. It allows us to plot graphs and charts using simpler syntax and makes it easier to adjust styles and colours. Most conveniently, since it is based on **Matplotlib**, many **Matplotlib** functions can be used to fine-tune **Seaborn** plots.

We import **NumPy**, **pandas**, **Matplotlib** and **Seaborn** as `np` , `pd` , `plt` and `sns` respectively as they are conventional abbreviations that make it faster to access the

features of these libraries (ie. type less code). Only the `pyplot` module of **Matplotlib** is imported. `pyplot` is a module of **Matplotlib** that allows graphs to be plotted using a MATLAB-like interface, albeit with Python code.

Finally, `%matplotlib inline` is executed to make our graphs and charts appear properly in our Jupyter notebook.

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

%matplotlib inline
```

Wrangling the data

Life expectancy

Let us use **pandas** to read `LifeExpectancyData-v2.csv` into a DataFrame called `life_exp`. The **pandas** function to read CSV files is `read_csv`. Since the CSV file is in the folder called `data`, and this notebook file is in the same directory as the folder, the path of the CSV file is `data/LifeExpectancyData-v2.csv`.

`life_exp.info()` gives us some information about the data, including the names of columns.

```
In [2]: life_exp = pd.read_csv('data/LifeExpectancyData-v2.csv')

life_exp.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2938 entries, 0 to 2937
Data columns (total 15 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   country                              2938 non-null   object
 1   Year                                  2938 non-null   int64
 2   Status                               2938 non-null   object
 3   Life expectancy                      2930 non-null   float64
 4   infant deaths                       2938 non-null   int64
 5   Adult Mortality                     2928 non-null   float64
 6   BMI                                  2904 non-null   float64
 7   Alcohol consumption                 2744 non-null   float64
 8   Hepatitis B                         2385 non-null   float64
 9   Measles                             2938 non-null   int64
10  Polio                               2919 non-null   float64
11  Diphtheria                          2919 non-null   float64
12  HIV/AIDS                           2938 non-null   float64
13  Income composition of resources     2771 non-null   float64
14  Schooling                           2775 non-null   float64
dtypes: float64(10), int64(3), object(2)
memory usage: 344.4+ KB
```

Some of the column labels have leading and trailing whitespaces, so we remove them using `strip()`. The stripped labels can then be applied through every column via a for loop that executes the `rename` method with every iteration. `inplace` is `True`

for the changes to be permanent. Additionally, the DataFrame is sorted using `sort_values` by `country`, then by `Year`. This gives us the same DataFrame but with the rows in ascending order of year.

Next, we copy this DataFrame into another variable called `life_exp_ori`. We will be making changes to `life_exp` as we continue to wrangle the data, so it is a good idea to keep a copy of the original DataFrame.

```
In [3]: for col in life_exp:
        life_exp.rename(columns = {col: col.strip()}, inplace = True)

        life_exp = life_exp.sort_values(['country', 'Year'])

        life_exp_ori = life_exp

        life_exp.head()
```

Out[3]:

	country	Year	Status	Life expectancy	infant deaths	Adult Mortality	BMI	Alcohol consumption	Hepat
15	Afghanistan	2000	Developing	54.8	88	321.0	12.2	0.01	6
14	Afghanistan	2001	Developing	55.3	88	316.0	12.6	0.01	6
13	Afghanistan	2002	Developing	56.2	88	3.0	13.0	0.01	6
12	Afghanistan	2003	Developing	56.7	87	295.0	13.4	0.01	6
11	Afghanistan	2004	Developing	57.0	87	293.0	13.8	0.02	6

Now we have a DataFrame that displays detailed information about life expectancies in various countries, with each country having separate entries for different years.

There is also data for various health and sociocultural factors that contribute to mortality rate. This amount of data is a little overwhelming, so some skimming is necessary.

Since `Life expectancy` is a responding variable of infant deaths and adult mortality, we can drop the two latter columns. `BMI` is influenced by health factors as well so we can keep that and drop the columns for diseases and alcohol consumption.

This should give us a subset of the original DataFrame containing the most essential data.

```
In [4]: life_exp = life_exp.filter(['country', 'Status', 'Life expectancy', 'BMI',
                                   'Income composition of resources', 'Schooling'])

        life_exp.head()
```

Out [4]:

	country	Status	Life expectancy	BMI	Income composition of resources	Schooling
15	Afghanistan	Developing	54.8	12.2	0.338	5.5
14	Afghanistan	Developing	55.3	12.6	0.340	5.9
13	Afghanistan	Developing	56.2	13.0	0.341	6.2
12	Afghanistan	Developing	56.7	13.4	0.373	6.5
11	Afghanistan	Developing	57.0	13.8	0.381	6.8

For this assignment, we will only be analysing the data for Southeast Asian countries (including East Timor). Using the `unique` function on a Series returns all the unique country names in this dataset, formatted as a **NumPy** array.

```
In [5]: life_exp['country'].unique()
```

```
Out[5]: array(['Afghanistan', 'Albania', 'Algeria', 'Angola',
      'Antigua and Barbuda', 'Argentina', 'Armenia', 'Australia',
      'Austria', 'Azerbaijan', 'Bahamas', 'Bahrain', 'Bangladesh',
      'Barbados', 'Belarus', 'Belgium', 'Belize', 'Benin', 'Bhutan',
      'Bolivia (Plurinational State of)', 'Bosnia and Herzegovina',
      'Botswana', 'Brazil', 'Brunei Darussalam', 'Bulgaria',
      'Burkina Faso', 'Burundi', 'Cabo Verde', 'Cambodia', 'Cameroon',
      'Canada', 'Central African Republic', 'Chad', 'Chile', 'China',
      'Colombia', 'Comoros', 'Congo', 'Cook Islands', 'Costa Rica',
      'Croatia', 'Cuba', 'Cyprus', 'Czechia', 'Côte d'Ivoire',
      'Democratic People's Republic of Korea',
      'Democratic Republic of the Congo', 'Denmark', 'Djibouti',
      'Dominica', 'Dominican Republic', 'Ecuador', 'Egypt',
      'El Salvador', 'Equatorial Guinea', 'Eritrea', 'Estonia',
      'Ethiopia', 'Fiji', 'Finland', 'France', 'Gabon', 'Gambia',
      'Georgia', 'Germany', 'Ghana', 'Greece', 'Grenada', 'Guatemala',
      'Guinea', 'Guinea-Bissau', 'Guyana', 'Haiti', 'Honduras',
      'Hungary', 'Iceland', 'India', 'Indonesia',
      'Iran (Islamic Republic of)', 'Iraq', 'Ireland', 'Israel', 'Italy',
      'Jamaica', 'Japan', 'Jordan', 'Kazakhstan', 'Kenya', 'Kiribati',
      'Kuwait', 'Kyrgyzstan', 'Lao People's Democratic Republic',
      'Latvia', 'Lebanon', 'Lesotho', 'Liberia', 'Libya', 'Lithuania',
      'Luxembourg', 'Madagascar', 'Malawi', 'Malaysia', 'Maldives',
      'Mali', 'Malta', 'Marshall Islands', 'Mauritania', 'Mauritius',
      'Mexico', 'Micronesia (Federated States of)', 'Monaco', 'Mongolia',
      'Montenegro', 'Morocco', 'Mozambique', 'Myanmar', 'Namibia',
      'Nauru', 'Nepal', 'Netherlands', 'New Zealand', 'Nicaragua',
      'Niger', 'Nigeria', 'Niue', 'Norway', 'Oman', 'Pakistan', 'Palau',
      'Panama', 'Papua New Guinea', 'Paraguay', 'Peru', 'Philippines',
      'Poland', 'Portugal', 'Qatar', 'Republic of Korea',
      'Republic of Moldova', 'Romania', 'Russian Federation', 'Rwanda',
      'Saint Kitts and Nevis', 'Saint Lucia',
      'Saint Vincent and the Grenadines', 'Samoa', 'San Marino',
      'Sao Tome and Principe', 'Saudi Arabia', 'Senegal', 'Serbia',
      'Seychelles', 'Sierra Leone', 'Singapore', 'Slovakia', 'Slovenia',
      'Solomon Islands', 'Somalia', 'South Africa', 'South Sudan',
      'Spain', 'Sri Lanka', 'Sudan', 'Suriname', 'Swaziland', 'Sweden',
      'Switzerland', 'Syrian Arab Republic', 'Tajikistan', 'Thailand',
      'The former Yugoslav republic of Macedonia', 'Timor-Leste', 'Togo',
      'Tonga', 'Trinidad and Tobago', 'Tunisia', 'Turkey',
      'Turkmenistan', 'Tuvalu', 'Uganda', 'Ukraine',
      'United Arab Emirates',
      'United Kingdom of Great Britain and Northern Ireland',
      'United Republic of Tanzania', 'United States of America',
      'Uruguay', 'Uzbekistan', 'Vanuatu',
      'Venezuela (Bolivarian Republic of)', 'Viet Nam', 'Yemen',
      'Zambia', 'Zimbabwe'], dtype=object)
```

The names are already arranged neatly in alphabetical order. Referring to [this table](#), we can easily take note of the names of Southeast Asian countries.

To make our data more viewable, some of the countries' names can be shortened to more conventional versions. We can do this by creating a dictionary mapping each name to its shortened version, then applying it into the `country` column using `replace`.

```
In [6]: renames = {'Brunei Darussalam': 'Brunei',
      'Timor-Leste': 'East Timor',
      'Lao People\'s Democratic Republic': 'Laos',
      'Viet Nam': 'Vietnam'
      }
```

```
life_exp['country'] = life_exp['country'].replace(renames)
```

Then, we can create a list containing all Southeast Asian countries including East Timor and obtain a subset of the original DataFrame containing only those countries via the `isin` method, that accepts our list as an argument.

```
In [7]: sea = ['Brunei', 'Cambodia', 'East Timor', 'Indonesia', 'Laos', 'Malaysia',
life_exp_sea = life_exp[life_exp['country'].isin(sea)]
life_exp_sea
```

```
Out[7]:
```

	country	Status	Life expectancy	BMI	Income composition of resources	Schooling
383	Brunei	Developing	74.4	26.1	0.818	13.4
382	Brunei	Developing	74.7	27.0	0.819	13.4
381	Brunei	Developing	74.8	28.0	0.820	13.3
380	Brunei	Developing	76.0	29.1	0.823	13.4
379	Brunei	Developing	76.4	3.1	0.828	13.7
...
2878	Vietnam	Developing	75.4	14.7	0.655	12.0
2877	Vietnam	Developing	75.6	15.3	0.662	12.2
2876	Vietnam	Developing	75.7	16.0	0.668	12.3
2875	Vietnam	Developing	75.9	16.7	0.675	12.5
2874	Vietnam	Developing	76.0	17.5	0.678	12.6

176 rows × 6 columns

We should now audit the data for missing or erroneous values, using `info` and `describe`.

```
In [8]: life_exp_sea.info()
life_exp_sea.describe()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 176 entries, 383 to 2874
Data columns (total 6 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   country                               176 non-null    object
1   Status                                176 non-null    object
2   Life expectancy                       176 non-null    float64
3   BMI                                   176 non-null    float64
4   Income composition of resources       176 non-null    float64
5   Schooling                             176 non-null    float64
dtypes: float64(4), object(2)
memory usage: 9.6+ KB
```

Out [8]:

	Life expectancy	BMI	Income composition of resources	Schooling
count	176.000000	176.000000	176.000000	176.000000
mean	70.035227	19.828977	0.643858	11.454545
std	6.220933	9.502169	0.140019	2.111325
min	57.700000	1.000000	0.000000	0.000000
25%	65.675000	13.975000	0.540000	10.475000
50%	68.300000	18.300000	0.641000	11.600000
75%	74.600000	27.000000	0.733250	12.725000
max	87.000000	41.200000	0.924000	15.400000

The output of `info` shows 176 non-null objects, which means there are no missing values. However, we deduce from the output of `describe` that there are discrepancies in the data for `BMI`, as evident from its minimum value of 1 (which is a realistically impossible value!).

One way to approach this is to impute the erroneous data with the mean of the correct data. The strategy is to replace the erroneous values with `NaN` and then set all `NaN` values with the mean using conditional slicing. The intermediate step of replacing these values with `NaN` is to ensure that the calculation of the mean does not include those values.

The lowest `BMI` value that is not obviously erroneous can be found in Vietnam's data ('not obviously' meaning the value fits into the trend of the other values), as shown below:

```
In [9]: life_exp_sea[life_exp_sea['country'] == 'Vietnam']
```

Out[9]:

	country	Status	Life expectancy	BMI	Income composition of resources	Schooling
2889	Vietnam	Developing	73.4	9.2	0.569	10.4
2888	Vietnam	Developing	73.6	9.6	0.576	10.6
2887	Vietnam	Developing	73.8	1.0	0.584	10.7
2886	Vietnam	Developing	74.0	1.4	0.592	10.9
2885	Vietnam	Developing	74.2	1.9	0.601	11.0
2884	Vietnam	Developing	74.4	11.3	0.609	11.1
2883	Vietnam	Developing	74.6	11.8	0.618	11.3
2882	Vietnam	Developing	74.7	12.3	0.625	11.4
2881	Vietnam	Developing	74.9	12.9	0.633	11.6
2880	Vietnam	Developing	75.0	13.4	0.641	11.7
2879	Vietnam	Developing	75.2	14.0	0.647	11.9
2878	Vietnam	Developing	75.4	14.7	0.655	12.0
2877	Vietnam	Developing	75.6	15.3	0.662	12.2
2876	Vietnam	Developing	75.7	16.0	0.668	12.3
2875	Vietnam	Developing	75.9	16.7	0.675	12.5
2874	Vietnam	Developing	76.0	17.5	0.678	12.6

Thus, we can define an erroneous BMI value as one below 9.2.

Using `loc`, we can slice (by index) all rows where the value of BMI is less than 9.2 and (by column) the BMI column. Then, using a simple assignment operator, we can set all these values to `np.nan`, which is a **NumPy** constant.

```
In [10]: life_exp_sea.loc[life_exp_sea['BMI'] < 9.2, 'BMI'] = np.nan
life_exp_sea
```


Out[10]:

	country	Status	Life expectancy	BMI	Income composition of resources	Schooling
383	Brunei	Developing	74.4	26.1	0.818	13.4
382	Brunei	Developing	74.7	27.0	0.819	13.4
381	Brunei	Developing	74.8	28.0	0.820	13.3
380	Brunei	Developing	76.0	29.1	0.823	13.4
379	Brunei	Developing	76.4	NaN	0.828	13.7
...
2878	Vietnam	Developing	75.4	14.7	0.655	12.0
2877	Vietnam	Developing	75.6	15.3	0.662	12.2
2876	Vietnam	Developing	75.7	16.0	0.668	12.3
2875	Vietnam	Developing	75.9	16.7	0.675	12.5
2874	Vietnam	Developing	76.0	17.5	0.678	12.6

176 rows × 6 columns

`life_exp_sea['country'].unique()` returns an array of the country names.

Using a for loop, we iterate through each country, calculating the mean BMI value for each country (**pandas's** `mean` method has an argument called `skipna` that is `True` by default, meaning it automatically excludes `NaN` values from calculation).

Then, we replace all `NaN` values for that country with the mean using `loc` and assignment. The by-index slice is all rows with the matching country name **and** where the BMI value is `NaN`, which can be easily identified via the `isnull` method. The by-column slice is the BMI column.

```
In [11]: for c in life_exp_sea['country'].unique():
          mean_bmi = life_exp_sea[life_exp_sea['country'] == c]['BMI'].mean()
          life_exp_sea.loc[(life_exp_sea['country'] == c) & (life_exp_sea['BMI'].isnull()), 'BMI'] = mean_bmi
          life_exp_sea
```

Out[11]:

	country	Status	Life expectancy	BMI	Income composition of resources	Schooling
383	Brunei	Developing	74.4	26.100000	0.818	13.4
382	Brunei	Developing	74.7	27.000000	0.819	13.4
381	Brunei	Developing	74.8	28.000000	0.820	13.3
380	Brunei	Developing	76.0	29.100000	0.823	13.4
379	Brunei	Developing	76.4	33.442857	0.828	13.7
...
2878	Vietnam	Developing	75.4	14.700000	0.655	12.0
2877	Vietnam	Developing	75.6	15.300000	0.662	12.2
2876	Vietnam	Developing	75.7	16.000000	0.668	12.3
2875	Vietnam	Developing	75.9	16.700000	0.675	12.5
2874	Vietnam	Developing	76.0	17.500000	0.678	12.6

176 rows × 6 columns

In addition, the values for `Income composition of resources` and `Schooling` for East Timor are 0 for the year 2000.

```
In [12]: life_exp_sea[life_exp_sea['country'] == 'East Timor']
```

Out[12]:

	country	Status	Life expectancy	BMI	Income composition of resources	Schooling
2616	East Timor	Developing	58.7	11.9	0.000	0.0
2615	East Timor	Developing	59.4	12.3	0.470	9.8
2614	East Timor	Developing	62.0	12.6	0.475	9.8
2613	East Timor	Developing	61.0	12.9	0.485	9.8
2612	East Timor	Developing	62.3	13.2	0.484	10.2
2611	East Timor	Developing	63.7	13.5	0.492	10.6
2610	East Timor	Developing	64.9	13.9	0.511	11.0
2609	East Timor	Developing	65.8	14.2	0.541	11.3
2608	East Timor	Developing	66.2	14.7	0.566	11.7
2607	East Timor	Developing	66.6	15.1	0.599	12.1
2606	East Timor	Developing	66.9	15.5	0.599	12.4
2605	East Timor	Developing	67.2	15.8	0.607	12.5
2604	East Timor	Developing	67.4	16.2	0.618	12.5
2603	East Timor	Developing	67.7	16.6	0.620	12.5
2602	East Timor	Developing	68.0	17.0	0.612	12.5
2601	East Timor	Developing	68.3	17.4	0.603	12.5

Since this is the first year in the dataset, it is not a good idea to simply impute these values with the mean. A simple look at the DataFrame and we can see that `Income composition of resources` grows at a consistent rate of 0.005 per year whereas `Schooling` remains at a constant value of 9.8 for the first few years. Hence, we can impute the zero values with dummy values of 0.465 and 9.8 respectively.

The `at` property allows quick access of the value in a specific cell of a DataFrame. We can then use assignment to directly replace the value.

```
In [13]: life_exp_sea.at[2616, 'Income composition of resources'] = 0.465
life_exp_sea.at[2616, 'Schooling'] = 9.8
```

```
life_exp_sea[life_exp_sea['country'] == 'East Timor'].head()
```

Out[13]:

	country	Status	Life expectancy	BMI	Income composition of resources	Schooling
2616	East Timor	Developing	58.7	11.9	0.465	9.8
2615	East Timor	Developing	59.4	12.3	0.470	9.8
2614	East Timor	Developing	62.0	12.6	0.475	9.8
2613	East Timor	Developing	61.0	12.9	0.485	9.8
2612	East Timor	Developing	62.3	13.2	0.484	10.2

Now we have a DataFrame that covers all the core data for Southeast Asian countries. This is useful for observing trends and growth of the data.

But what if we need a summary? Each country has many entries in the DataFrame, one for each year in the dataset. We can go one step further and obtain aggregated values for each statistic. More specifically, we reorganise the DataFrame so that it shows only the mean values for each statistic. Mean is a good measure of central tendency as its calculation involves every value in the dataset. As the responding variable, life expectancy will also get a column for its max values as well, so that they may be compared with the corresponding mean values.

Since `country` is the set of values we want to separate our data by, we will utilise the `groupby` function to group the data by `country`.

We first create a dictionary of keyword arguments to be passed into the `agg` function. For each entry in the dictionary, the string before the colon is the name of the column to be added. The tuple after the colon contains first the label of the column whose data is to be aggregated, then the function used for aggregation.

The `groupby` object is then assigned to a new variable called `life_exp_sea_agg`. `as_index` is set to `False` to prevent the country column from being used as the index of the DataFrame. `agg` is then applied to this DataFrame with the keyword arguments. The characters `**` must be included to denote that our argument is a dictionary of keyword arguments to be unpacked.

```
In [14]: fun = {'max_life_expectancy': ('Life expectancy', 'max'),
               'mean_life_expectancy': ('Life expectancy', 'mean'),
               'mean_BMI': ('BMI', 'mean'),
               'mean_income_composition_of_resources': ('Income composition of resources', 'mean'),
               'mean_schooling': ('Schooling', 'mean'),
               }

life_exp_sea_agg = life_exp_sea.groupby(['country'], as_index = False).agg(**fun)

life_exp_sea_agg
```

Out[14]:

	country	max_life_expectancy	mean_life_expectancy	mean_BMI	mean_income_compos
0	Brunei	78.3	76.48750	33.442857	
1	Cambodia	68.7	64.34375	15.362500	
2	East Timor	68.3	64.75625	14.550000	
3	Indonesia	69.1	67.55625	21.120000	
4	Laos	65.7	62.38125	16.057143	
5	Malaysia	75.0	73.75625	32.742857	
6	Myanmar	66.6	64.20000	18.100000	
7	Philippines	68.5	67.57500	21.592857	
8	Singapore	87.0	81.47500	31.069231	
9	Thailand	74.9	73.08125	25.946154	
10	Vietnam	76.0	74.77500	13.438462	

It is a good idea to include `Status` as it can be linked to the numerical statistics in our new DataFrame.

First, we want to see if any countries changed their status throughout the years. We slice the `life_exp_sea` DataFrame using `iloc[:, 0:2]` means include all rows and include only the first two columns.

In [15]: `life_exp_sea.iloc[:, 0:2]`

Out[15]:

	country	Status
383	Brunei	Developing
382	Brunei	Developing
381	Brunei	Developing
380	Brunei	Developing
379	Brunei	Developing
...
2878	Vietnam	Developing
2877	Vietnam	Developing
2876	Vietnam	Developing
2875	Vietnam	Developing
2874	Vietnam	Developing

176 rows × 2 columns

Then, we remove duplicate rows using `drop_duplicates`. If any countries changed their status, they will appear twice in the DataFrame; once with `Developing` status,

and once with `Developed` status. Otherwise, there will be as many rows as there are countries.

We also use `sort_values` to reorder the DataFrame by alphabetical order of country names. This is necessary as `Timor-Leste` has been renamed to `East Timor`.

```
In [16]: life_exp_sea_status = life_exp_sea.iloc[:, 0:2].drop_duplicates().sort_values(
life_exp_sea_status
```

```
Out[16]:
```

	country	Status
383	Brunei	Developing
479	Cambodia	Developing
2616	East Timor	Developing
1217	Indonesia	Developing
1441	Laos	Developing
1601	Malaysia	Developing
1795	Myanmar	Developing
2038	Philippines	Developing
2328	Singapore	Developed
2584	Thailand	Developing
2889	Vietnam	Developing

So, for the years covered in this dataset, no countries experienced a change in development status. We do not need to worry about this influencing our data. We can now proceed to add the other columns from `life_exp_sea_agg` (using `merge`) to `life_sea_exp_status`, then reassign this complete DataFrame to the `life_exp_sea_agg` variable. We do not need to pass in any arguments as the `country` column is found in both DataFrames, so **pandas** uses that as the default column to merge on.

```
In [17]: life_exp_sea_agg = life_exp_sea_status.merge(life_exp_sea_agg)
life_exp_sea_agg
```

Out[17]:

	country	Status	max_life_expectancy	mean_life_expectancy	mean_BMI	mean_income
0	Brunei	Developing	78.3	76.48750	33.442857	
1	Cambodia	Developing	68.7	64.34375	15.362500	
2	East Timor	Developing	68.3	64.75625	14.550000	
3	Indonesia	Developing	69.1	67.55625	21.120000	
4	Laos	Developing	65.7	62.38125	16.057143	
5	Malaysia	Developing	75.0	73.75625	32.742857	
6	Myanmar	Developing	66.6	64.20000	18.100000	
7	Philippines	Developing	68.5	67.57500	21.592857	
8	Singapore	Developed	87.0	81.47500	31.069231	
9	Thailand	Developing	74.9	73.08125	25.946154	
10	Vietnam	Developing	76.0	74.77500	13.438462	

Population

Now, let us integrate data from other datasets into our DataFrame. Population is a good choice as we can observe how population influences the data we currently have. We import `2020-Population.csv`, a dataset containing populations of countries from 1950 to 2020.

```
In [18]: population = pd.read_csv('data/2020-Population.csv')
population.info()
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 305 entries, 0 to 304

Data columns (total 78 columns):

#	Column	Non-Null Count	Dtype
0	Unnamed: 0	299 non-null	object
1	Unnamed: 1	290 non-null	object
2	Unnamed: 2	290 non-null	object
3	Unnamed: 3	83 non-null	object
4	Unnamed: 4	290 non-null	object
5	Unnamed: 5	290 non-null	object
6	Unnamed: 6	290 non-null	object
7	Unnamed: 7	291 non-null	object
8	Unnamed: 8	290 non-null	object
9	Unnamed: 9	290 non-null	object
10	Unnamed: 10	290 non-null	object
11	Unnamed: 11	290 non-null	object
12	Unnamed: 12	290 non-null	object
13	Unnamed: 13	290 non-null	object
14	Unnamed: 14	290 non-null	object
15	Unnamed: 15	290 non-null	object
16	Unnamed: 16	290 non-null	object
17	Unnamed: 17	290 non-null	object
18	Unnamed: 18	290 non-null	object
19	Unnamed: 19	290 non-null	object
20	Unnamed: 20	290 non-null	object
21	Unnamed: 21	290 non-null	object
22	Unnamed: 22	290 non-null	object
23	Unnamed: 23	290 non-null	object
24	Unnamed: 24	290 non-null	object
25	Unnamed: 25	290 non-null	object
26	Unnamed: 26	290 non-null	object
27	Unnamed: 27	290 non-null	object
28	Unnamed: 28	290 non-null	object
29	Unnamed: 29	290 non-null	object
30	Unnamed: 30	290 non-null	object
31	Unnamed: 31	290 non-null	object
32	Unnamed: 32	290 non-null	object
33	Unnamed: 33	290 non-null	object
34	Unnamed: 34	290 non-null	object
35	Unnamed: 35	290 non-null	object
36	Unnamed: 36	290 non-null	object
37	Unnamed: 37	290 non-null	object
38	Unnamed: 38	290 non-null	object
39	Unnamed: 39	290 non-null	object
40	Unnamed: 40	290 non-null	object
41	Unnamed: 41	290 non-null	object
42	Unnamed: 42	290 non-null	object
43	Unnamed: 43	290 non-null	object
44	Unnamed: 44	290 non-null	object
45	Unnamed: 45	290 non-null	object
46	Unnamed: 46	290 non-null	object
47	Unnamed: 47	290 non-null	object
48	Unnamed: 48	290 non-null	object
49	Unnamed: 49	290 non-null	object
50	Unnamed: 50	290 non-null	object
51	Unnamed: 51	290 non-null	object
52	Unnamed: 52	290 non-null	object
53	Unnamed: 53	290 non-null	object
54	Unnamed: 54	290 non-null	object
55	Unnamed: 55	290 non-null	object
56	Unnamed: 56	290 non-null	object
57	Unnamed: 57	290 non-null	object
58	Unnamed: 58	290 non-null	object


```

59 Unnamed: 59 290 non-null object
60 Unnamed: 60 290 non-null object
61 Unnamed: 61 290 non-null object
62 Unnamed: 62 290 non-null object
63 Unnamed: 63 290 non-null object
64 Unnamed: 64 290 non-null object
65 Unnamed: 65 290 non-null object
66 Unnamed: 66 290 non-null object
67 Unnamed: 67 290 non-null object
68 Unnamed: 68 290 non-null object
69 Unnamed: 69 290 non-null object
70 Unnamed: 70 290 non-null object
71 Unnamed: 71 290 non-null object
72 Unnamed: 72 290 non-null object
73 Unnamed: 73 290 non-null object
74 Unnamed: 74 290 non-null object
75 Unnamed: 75 290 non-null object
76 Unnamed: 76 290 non-null object
77 Unnamed: 77 290 non-null object
dtypes: object(78)
memory usage: 186.0+ KB

```

The dataset looks to be messy. We will have to wrangle it. First, we look at its head.

```
In [19]: population.head()
```

```

Out[19]:      Unnamed: 0  Unnamed: 1  Unnamed: 2  Unnamed: 3  Unnamed: 4  Unnamed: 5  Unnamed: 6  Unnamed: 7
0      NaN      NaN      NaN      NaN      NaN      NaN      NaN      NaN
1      NaN      NaN      NaN      NaN      NaN      NaN      NaN      NaN
2      NaN      NaN      NaN      NaN      NaN      NaN      NaN      NaN
3  United Nations      NaN      NaN      NaN      NaN      NaN      NaN      NaN
4  Population Division      NaN      NaN      NaN      NaN      NaN      NaN      NaN

```

5 rows × 78 columns

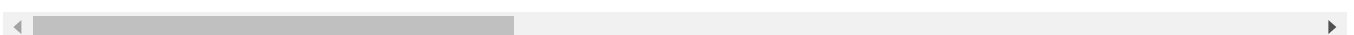
The head of the DataFrame does not tell us much about the content of the dataset. Let us look at the first 20 rows instead.

```
In [20]: population[:20]
```

Out[20]:

	Unnamed: 0	Unnamed: 1	Unnamed: 2	Unnamed: 3	Unnamed: 4	Unnar
0	NaN	NaN	NaN	NaN	NaN	
1	NaN	NaN	NaN	NaN	NaN	
2	NaN	NaN	NaN	NaN	NaN	
3	United Nations	NaN	NaN	NaN	NaN	
4	Population Division	NaN	NaN	NaN	NaN	
5	Department of Economic and Social Affairs	NaN	NaN	NaN	NaN	
6	NaN	NaN	NaN	NaN	NaN	
7	World Population Prospects 2019	NaN	NaN	NaN	NaN	
8	File POP/1-1: Total population (both sexes com...	NaN	NaN	NaN	NaN	
9	Estimates, 1950 - 2020	NaN	NaN	NaN	NaN	
10	POP/DB/WPP/Rev.2019/POP/F01-1	NaN	NaN	NaN	NaN	
11	© August 2019 by United Nations, made availabl...	NaN	NaN	NaN	NaN	
12	Suggested citation: United Nations, Department...	NaN	NaN	NaN	NaN	
13	NaN	NaN	NaN	NaN	NaN	
14	NaN	NaN	NaN	NaN	NaN	
15	Index	Variant	Region, subregion, country or area *	Notes	Country code	
16	1	Estimates	WORLD	NaN	900	
17	2	Estimates	UN development groups	a	1803	Label/Sep
18	3	Estimates	More developed regions	b	901	Develo
19	4	Estimates	Less developed regions	c	902	Develo

20 rows × 78 columns



Thus, we can see the data really only begins starting from row 16, with row 15 being

the headers. All the rows above it just contain information about the dataset that we can safely drop.

```
In [21]: population = population.drop(index = range(15))
population.head()
```

```
Out[21]:
```

	Unnamed: 0	Unnamed: 1	Unnamed: 2	Unnamed: 3	Unnamed: 4	Unnamed: 5	Unnamed: 6
15	Index	Variant	Region, subregion, country or area *	Notes	Country code	Type	Parent code
16	1	Estimates	World	NaN	900	World	0
17	2	Estimates	UN development groups	a	1803	Label/Separator	900
18	3	Estimates	More developed regions	b	901	Development Group	1803
19	4	Estimates	Less developed regions	c	902	Development Group	1803

5 rows × 78 columns

Now, we set the first row of the DataFrame as the headers of the columns. The original row can then be removed from the DataFrame. The column `Index` can also be removed as we will be working with the built-in indexes of the DataFrame.

Since we used row 15 as our column headers, the index of our DataFrame will now have a name called '15'. We remove it using `rename_axis`.

```
In [22]: population.columns = population.iloc[0]
population = population.drop(index = 15, columns = 'Index')
population.rename_axis(None, axis = 1, inplace = True)
population.head()
```

Out [22] :

	Variant	Region, subregion, country or area *	Notes	Country code	Type	Parent code	1950	1951	1952	1953
16	Estimates	WORLD	NaN	900	World	0	2 536 431	2 584 034	2 630 862	2 677 609
17	Estimates	UN development groups	a	1803	Label/Separator	900
18	Estimates	More developed regions	b	901	Development Group	1803	814 819	824 004	833 720	843 788
19	Estimates	Less developed regions	c	902	Development Group	1803	1 721 612	1 760 031	1 797 142	1 833 822
20	Estimates	Least developed countries	d	941	Development Group	902	195 428	199 180	203 015	206 986

5 rows × 77 columns



Again, we only want the data for Southeast Asian countries. We can reuse the code we used above, but let us first rename the label of the column of country/region names to something shorter for convenience. We will only work with countries in the end anyways.

```
In [23]: population.rename(columns = {'Region, subregion, country or area *': 'country'})
population.head()
```

Out [23]:

	Variant	country	Notes	Country code	Type	Parent code	1950	1951	1952	1953
16	Estimates	WORLD	NaN	900	World	0	536 431	584 034	630 862	677 609
17	Estimates	UN development groups	a	1803	Label/Separator	900
18	Estimates	More developed regions	b	901	Development Group	1803	814 819	824 004	833 720	843 788
19	Estimates	Less developed regions	c	902	Development Group	1803	721 612	760 031	797 142	833 822
20	Estimates	Least developed countries	d	941	Development Group	902	195 428	199 180	203 015	206 986

5 rows × 77 columns



We use `unique` again to list down the names of countries in an array. We are trying to note the names of Southeast Asian countries and see if they match the conventional versions we used before. Since the names are not in alphabetical order, we can use `np.sort`, which sorts a **NumPy** array and returns a copy.

In [24]: `np.sort(population['country'].unique())`

```

Out[24]: array(['AUSTRALIA/NEW ZEALAND', 'Afghanistan', 'Africa', 'Albania',
                'Algeria', 'American Samoa', 'Andorra', 'Angola', 'Anguilla',
                'Antigua and Barbuda', 'Argentina', 'Armenia', 'Aruba', 'Asia',
                'Australia', 'Austria', 'Azerbaijan', 'Bahamas', 'Bahrain',
                'Bangladesh', 'Barbados', 'Belarus', 'Belgium', 'Belize', 'Benin',
                'Bermuda', 'Bhutan', 'Bolivia (Plurinational State of)',
                'Bonaire, Sint Eustatius and Saba', 'Bosnia and Herzegovina',
                'Botswana', 'Brazil', 'British Virgin Islands',
                'Brunei Darussalam', 'Bulgaria', 'Burkina Faso', 'Burundi',
                'CENTRAL AND SOUTHERN ASIA', 'Cabo Verde', 'Cambodia', 'Cameroon',
                'Canada', 'Caribbean', 'Cayman Islands',
                'Central African Republic', 'Central America', 'Central Asia',
                'Chad', 'Channel Islands', 'Chile', 'China',
                'China, Hong Kong SAR', 'China, Macao SAR',
                'China, Taiwan Province of China', 'Colombia', 'Comoros', 'Congo',
                'Cook Islands', 'Costa Rica', 'Croatia', 'Cuba', 'Curaçao',
                'Cyprus', 'Czechia', 'Côte d'Ivoire',
                'Dem. People's Republic of Korea',
                'Democratic Republic of the Congo', 'Denmark', 'Djibouti',
                'Dominica', 'Dominican Republic', 'EASTERN AND SOUTH-EASTERN ASI
A',
                'EUROPE', 'EUROPE AND NORTHERN AMERICA', 'Eastern Africa',
                'Eastern Asia', 'Eastern Europe', 'Ecuador', 'Egypt',
                'El Salvador', 'Equatorial Guinea', 'Eritrea', 'Estonia',
                'Eswatini', 'Ethiopia', 'Europe', 'Falkland Islands (Malvinas)',
                'Faroe Islands', 'Fiji', 'Finland', 'France', 'French Guiana',
                'French Polynesia', 'Gabon', 'Gambia', 'Geographic regions',
                'Georgia', 'Germany', 'Ghana', 'Gibraltar', 'Greece', 'Greenland',
                'Grenada', 'Guadeloupe', 'Guam', 'Guatemala', 'Guinea',
                'Guinea-Bissau', 'Guyana', 'Haiti', 'High-income countries',
                'Holy See', 'Honduras', 'Hungary', 'Iceland', 'India', 'Indonesi
a',
                'Iran (Islamic Republic of)', 'Iraq', 'Ireland', 'Isle of Man',
                'Israel', 'Italy', 'Jamaica', 'Japan', 'Jordan', 'Kazakhstan',
                'Kenya', 'Kiribati', 'Kuwait', 'Kyrgyzstan',
                'LATIN AMERICA AND THE CARIBBEAN',
                'Land-locked Developing Countries (LLDC)',
                'Lao People's Democratic Republic',
                'Latin America and the Caribbean', 'Latvia',
                'Least developed countries', 'Lebanon', 'Lesotho',
                'Less developed regions',
                'Less developed regions, excluding China',
                'Less developed regions, excluding least developed countries',
                'Liberia', 'Libya', 'Liechtenstein', 'Lithuania',
                'Low-income countries', 'Lower-middle-income countries',
                'Luxembourg', 'Madagascar', 'Malawi', 'Malaysia', 'Maldives',
                'Mali', 'Malta', 'Marshall Islands', 'Martinique', 'Mauritania',
                'Mauritius', 'Mayotte', 'Melanesia', 'Mexico', 'Micronesia',
                'Micronesia (Fed. States of)', 'Middle Africa',
                'Middle-income countries', 'Monaco', 'Mongolia', 'Montenegro',
                'Montserrat', 'More developed regions', 'Morocco', 'Mozambique',
                'Myanmar', 'NORTHERN AFRICA AND WESTERN ASIA', 'NORTHERN AMERICA',
                'Namibia', 'Nauru', 'Nepal', 'Netherlands', 'New Caledonia',
                'New Zealand', 'Nicaragua', 'Niger', 'Nigeria', 'Niue',
                'No income group available', 'North Macedonia', 'Northern Africa',
                'Northern America', 'Northern Europe', 'Northern Mariana Islands',
                'Norway', 'OCEANIA (EXCLUDING AUSTRALIA AND NEW ZEALAND)',
                'Oceania', 'Oman', 'Pakistan', 'Palau', 'Panama',
                'Papua New Guinea', 'Paraguay', 'Peru', 'Philippines', 'Poland',
                'Polynesia', 'Portugal', 'Puerto Rico', 'Qatar',
                'Republic of Korea', 'Republic of Moldova', 'Romania',
                'Russian Federation', 'Rwanda', 'Réunion', 'SUB-SAHARAN AFRICA',
                'Saint Barthélemy', 'Saint Helena', 'Saint Kitts and Nevis',
                'Saint Lucia', 'Saint Martin (French part)',

```

```
'Saint Pierre and Miquelon', 'Saint Vincent and the Grenadines',
'Samoa', 'San Marino', 'Sao Tome and Principe', 'Saudi Arabia',
'Senegal', 'Serbia', 'Seychelles', 'Sierra Leone', 'Singapore',
'Sint Maarten (Dutch part)', 'Slovakia', 'Slovenia',
'Small Island Developing States (SIDS)', 'Solomon Islands',
'Somalia', 'South Africa', 'South America', 'South Sudan',
'South-Eastern Asia', 'Southern Africa', 'Southern Asia',
'Southern Europe', 'Spain', 'Sri Lanka', 'State of Palestine',
'Sudan', 'Suriname', 'Sustainable Development Goal (SDG) regions',
'Sweden', 'Switzerland', 'Syrian Arab Republic', 'Tajikistan',
'Thailand', 'Timor-Leste', 'Togo', 'Tokelau', 'Tonga',
'Trinidad and Tobago', 'Tunisia', 'Turkey', 'Turkmenistan',
'Turks and Caicos Islands', 'Tuvalu', 'UN development groups',
'Uganda', 'Ukraine', 'United Arab Emirates', 'United Kingdom',
'United Republic of Tanzania', 'United States Virgin Islands',
'United States of America', 'Upper-middle-income countries',
'Uruguay', 'Uzbekistan', 'Vanuatu',
'Venezuela (Bolivarian Republic of)', 'Viet Nam', 'WORLD',
'Wallis and Futuna Islands', 'Western Africa', 'Western Asia',
'Western Europe', 'Western Sahara', 'World Bank income groups',
'Yemen', 'Zambia', 'Zimbabwe'], dtype=object)
```

Looks like the names for to this dataset have to be shortened as well. Fortunately, their 'long' versions are the same as the ones from `LifeExpectancyData-v2.csv`. We still have the `renames` dictionary as well as the list of Southeast Asian country names, `sea`, from earlier, which we can safely reuse.

```
renames = {'Brunei Darussalam': 'Brunei',
           'Timor-Leste': 'East Timor',
           'Lao People\'s Democratic Republic': 'Laos',
           'Viet Nam': 'Vietnam'
          }

sea = ['Brunei', 'Cambodia', 'East Timor', 'Indonesia',
       'Laos', 'Malaysia', 'Myanmar', 'Philippines', 'Singapore',
       'Thailand', 'Vietnam']
```

Just like before, the necessary country names are shortened using `replace` and `isin` is used to filter out the data for Southeast Asian countries.

For a clean appearance, we also use `reset_index` so the indexes start from 0. `drop` is set to `True` so there would not be a new column for the original indexes added to the DataFrame. The end result is assigned to a new variable called `population_sea`.

Finally, the DataFrame is sorted by `country` as `Timor-Leste` has been renamed to `East Timor`.

```
In [25]: population['country'] = population['country'].replace(renames)
population_sea = population[population['country'].isin(sea)].reset_index(drop=True)
population_sea = population_sea.sort_values('country')

population_sea
```

Out[25]:

	Variant	country	Notes	Country code	Type	Parent code	1950	1951	1952	1953	...
0	Estimates	Brunei	NaN	96	Country/Area	920	48	51	54	57	...
1	Estimates	Cambodia	NaN	116	Country/Area	920	4 433	4 538	4 656	4 783	...
9	Estimates	East Timor	NaN	626	Country/Area	920	415	419	423	428	...
2	Estimates	Indonesia	NaN	360	Country/Area	920	69 543	70 849	72 275	73 821	...
3	Estimates	Laos	NaN	418	Country/Area	920	1 683	1 723	1 764	1 806	...
4	Estimates	Malaysia	13	458	Country/Area	920	6 110	6 271	6 450	6 639	...
5	Estimates	Myanmar	NaN	104	Country/Area	920	17 780	18 104	18 441	18 793	...
6	Estimates	Philippines	NaN	608	Country/Area	920	18 580	19 247	19 945	20 670	...
7	Estimates	Singapore	NaN	702	Country/Area	920	1 022	1 068	1 120	1 178	...
8	Estimates	Thailand	NaN	764	Country/Area	920	20 710	21 263	21 838	22 437	...
10	Estimates	Vietnam	NaN	704	Country/Area	920	24 810	25 365	25 977	26 646	...

11 rows × 77 columns



We will be merging data for the gross domestic product (GDP) of each country later. The GDP data is from 2019. Hence, we just need the population values for 2019. Using `filter`, we extract only the country names and values for 2019 from `population_sea`.

```
In [26]: population_sea = population_sea.filter(['country', '2019'])
population_sea
```


Out [26]:

	country	2019
0	Brunei	433
1	Cambodia	16 487
9	East Timor	1 293
2	Indonesia	270 626
3	Laos	7 169
4	Malaysia	31 950
5	Myanmar	54 045
6	Philippines	108 117
7	Singapore	5 804
8	Thailand	69 626
10	Vietnam	96 462

Looks good. But the following output shows otherwise:

In [27]: `population_sea.dtypes`

Out[27]:

```
country    object
2019       object
dtype: object
```

The values of `2019` are strings! This will give us trouble if we attempt to visualise the data later.

Before we convert the strings to integers, we have to remove the whitespaces or Python would not be able to cast them. For this, we use the `apply` function with a lambda expression that replaces all whitespaces with empty strings. Then, the strings can be converted into integers by calling the `to_numeric` function on the `2019` series.

In addition, the values of `2019` are in thousands. This can be easily verified by searching for the exact population of a few of the countries using Google. We prefer to use exact values for our DataFrame, so we can multiply every value in `2019` by 1000.

In [28]:

```
population_sea['2019'] = population_sea['2019'].apply(lambda x: x.replace(' ', ''))
population_sea['2019'] = pd.to_numeric(population_sea['2019'])
population_sea['2019'] = population_sea['2019'] * 1000

population_sea
```

Out[28]:

	country	2019
0	Brunei	433000
1	Cambodia	16487000
9	East Timor	1293000
2	Indonesia	270626000
3	Laos	7169000
4	Malaysia	31950000
5	Myanmar	54045000
6	Philippines	108117000
7	Singapore	5804000
8	Thailand	69626000
10	Vietnam	96462000

Let us verify that the values have been casted to integers correctly.

In [29]: `population_sea.dtypes`

Out[29]:

country	object
2019	int64
dtype:	object

Now, we merge the DataFrame to `life_exp_sea_agg`, renaming the column `2019` to just `population`. We will reassign this newly combined DataFrame to a new variable called `life_exp_sea_pop`.

In [30]:

```
population_sea.rename(columns = {'2019': 'population'}, inplace = True)
life_exp_sea_pop = life_exp_sea_agg.merge(population_sea)

life_exp_sea_pop
```

Out[30]:

	country	Status	max_life_expectancy	mean_life_expectancy	mean_BMI	mean_income
0	Brunei	Developing	78.3	76.48750	33.442857	
1	Cambodia	Developing	68.7	64.34375	15.362500	
2	East Timor	Developing	68.3	64.75625	14.550000	
3	Indonesia	Developing	69.1	67.55625	21.120000	
4	Laos	Developing	65.7	62.38125	16.057143	
5	Malaysia	Developing	75.0	73.75625	32.742857	
6	Myanmar	Developing	66.6	64.20000	18.100000	
7	Philippines	Developing	68.5	67.57500	21.592857	
8	Singapore	Developed	87.0	81.47500	31.069231	
9	Thailand	Developing	74.9	73.08125	25.946154	
10	Vietnam	Developing	76.0	74.77500	13.438462	

GDP

We can now proceed to integrate the data from `2019-GDP.csv` into our DataFrame.

```
In [31]: gdp = pd.read_csv('data/2019-GDP.csv')
gdp.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 244 entries, 0 to 243
Data columns (total 6 columns):
 #   Column                                Non-Null Count  Dtype  
---  -
 0   Unnamed: 0                            229 non-null   object 
 1   Gross domestic product 2019          209 non-null   object 
 2   Unnamed: 2                            0 non-null     float64
 3   Unnamed: 3                            230 non-null   object 
 4   Unnamed: 4                            231 non-null   object 
 5   Unnamed: 5                            8 non-null     object 
dtypes: float64(1), object(5)
memory usage: 11.6+ KB
```

Again, this dataset will require wrangling from the look of its column labels. We first read its head.

```
In [32]: gdp.head()
```

```
Out[32]:
```

	Unnamed: 0	Gross domestic product 2019	Unnamed: 2	Unnamed: 3	Unnamed: 4	Unnamed: 5
0	NaN	NaN	NaN	NaN	NaN	NaN
1	NaN	NaN	NaN	NaN	(millions of	NaN
2	NaN	Ranking	NaN	Economy	US dollars)	NaN
3	NaN	NaN	NaN	NaN	NaN	NaN
4	USA	1	NaN	United States	21,427,700	NaN

Similar to the `2020-Population.csv` file, the head does not tell us much. We will need to look at more of the DataFrame and tidy it up.

```
In [33]: gdp.iloc[:20]
```

Out[33]:

	Unnamed: 0	Gross domestic product 2019	Unnamed: 2	Unnamed: 3	Unnamed: 4	Unnamed: 5
0	NaN	NaN	NaN	NaN	NaN	NaN
1	NaN	NaN	NaN	NaN	(millions of	NaN
2	NaN	Ranking	NaN	Economy	US dollars)	NaN
3	NaN	NaN	NaN	NaN	NaN	NaN
4	USA	1	NaN	United States	21,427,700	NaN
5	CHN	2	NaN	China	14,342,903	NaN
6	JPN	3	NaN	Japan	5,081,770	NaN
7	DEU	4	NaN	Germany	3,845,630	NaN
8	IND	5	NaN	India	2,875,142	NaN
9	GBR	6	NaN	United Kingdom	2,827,113	NaN
10	FRA	7	NaN	France	2,715,518	NaN
11	ITA	8	NaN	Italy	2,001,244	NaN
12	BRA	9	NaN	Brazil	1,839,758	NaN
13	CAN	10	NaN	Canada	1,736,426	NaN
14	RUS	11	NaN	Russian Federation	1,699,877	a
15	KOR	12	NaN	Korea, Rep.	1,642,383	NaN
16	ESP	13	NaN	Spain	1,394,116	NaN
17	AUS	14	NaN	Australia	1,392,681	NaN
18	MEX	15	NaN	Mexico	1,258,287	NaN
19	IDN	16	NaN	Indonesia	1,119,191	NaN

The columns labelled `Unnamed: 2` and `Unnamed: 5` consist mostly of `NaN` values. We will remove them along with rows 0, 1 and 3, which do not provide meaningful data as well.

```
In [34]: gdp = gdp.drop(index = [0, 1, 3], columns = ['Unnamed: 2', 'Unnamed: 5'])
gdp.head()
```

Out[34]:

	Unnamed: 0	Gross domestic product 2019	Unnamed: 3	Unnamed: 4
2	NaN	Ranking	Economy	US dollars)
4	USA	1	United States	21,427,700
5	CHN	2	China	14,342,903
6	JPN	3	Japan	5,081,770
7	DEU	4	Germany	3,845,630

Just like we observed in `2020-Population.csv`, the column headers of this DataFrame are in the first row. This row should be set as the column labels of this DataFrame. `rename_axis` is also used as we do not want the index to be labelled.

```
In [35]: gdp.columns = gdp.iloc[0]
gdp = gdp.drop(index = 2)
gdp.rename_axis(None, axis = 1, inplace = True)

gdp.head()
```

```
Out[35]:
```

	NaN	Ranking	Economy	US dollars)
4	USA	1	United States	21,427,700
5	CHN	2	China	14,342,903
6	JPN	3	Japan	5,081,770
7	DEU	4	Germany	3,845,630
8	IND	5	India	2,875,142

Now we have column headers, but they still do not make much sense (eg. `Economy` is not a correct label). Using `rename`, we give each column a more accurate name. Based on the original CSV file, the GDP is in millions of US dollars, so we have to mention that in the label as well.

Using `np.nan`, we can directly refer to the `NaN` value and rename it to `country code`. `'NaN': 'country code'` does not work in this case as `NaN` is not a string.

```
In [36]: gdp.rename(columns = {np.nan: 'country code',
                                'Economy': 'country',
                                'US dollars)': 'GDP (in millions USD)'},
                    inplace = True)

gdp.head()
```

```
Out[36]:
```

	country code	Ranking	country	GDP (in millions USD)
4	USA	1	United States	21,427,700
5	CHN	2	China	14,342,903
6	JPN	3	Japan	5,081,770
7	DEU	4	Germany	3,845,630
8	IND	5	India	2,875,142

To ensure the country names for this dataset match the conventional version, we call `unique` again.

```
In [37]: gdp['country'].unique()
```

```

Out[37]: array(['United States', 'China', 'Japan', 'Germany', 'India',
               'United Kingdom', 'France', 'Italy', 'Brazil', 'Canada',
               'Russian Federation', 'Korea, Rep.', 'Spain', 'Australia',
               'Mexico', 'Indonesia', 'Netherlands', 'Saudi Arabia', 'Turkey',
               'Switzerland', 'Poland', 'Thailand', 'Sweden', 'Belgium',
               'Argentina', 'Nigeria', 'Austria', 'Iran, Islamic Rep.',
               'United Arab Emirates', 'Norway', 'Israel', 'Ireland',
               'Philippines', 'Singapore', 'Hong Kong SAR, China', 'Malaysia',
               'South Africa', 'Denmark', 'Colombia', 'Egypt, Arab Rep.',
               'Bangladesh', 'Chile', 'Pakistan', 'Finland', 'Vietnam', 'Romani
a',
               'Czech Republic', 'Portugal', 'Iraq', 'Peru', 'Greece',
               'New Zealand', 'Qatar', 'Kazakhstan', 'Algeria', 'Hungary',
               'Ukraine', 'Kuwait', 'Morocco', 'Ecuador', 'Slovak Republic',
               'Puerto Rico', 'Cuba', 'Ethiopia', 'Kenya', 'Angola',
               'Dominican Republic', 'Sri Lanka', 'Oman', 'Guatemala', 'Myanmar',
               'Luxembourg', 'Bulgaria', 'Ghana', 'Panama', 'Tanzania', 'Belaru
s',
               'Costa Rica', 'Croatia', "Côte d'Ivoire", 'Uzbekistan', 'Uruguay',
               'Lithuania', 'Macao SAR, China', 'Slovenia', 'Lebanon', 'Libya',
               'Serbia', 'Azerbaijan', 'Congo, Dem. Rep.', 'Jordan', 'Bolivia',
               'Turkmenistan', 'Tunisia', 'Cameroon', 'Bahrain', 'Paraguay',
               'Uganda', 'Latvia', 'Estonia', 'Nepal', 'Yemen, Rep.', 'Cambodia',
               'El Salvador', 'Honduras', 'Papua New Guinea', 'Cyprus', 'Icelan
d',
               'Trinidad and Tobago', 'Senegal', 'Zambia', 'Zimbabwe',
               'Bosnia and Herzegovina', 'Afghanistan', 'Sudan', 'Botswana',
               'Lao PDR', 'Georgia', 'Mali', 'Gabon', 'Jamaica', 'Burkina Faso',
               'Albania', 'Mozambique', 'Malta', 'West Bank and Gaza', 'Benin',
               'Mauritius', 'Madagascar', 'Mongolia', 'Armenia', 'Guinea',
               'Brunei Darussalam', 'Niger', 'Bahamas, The', 'North Macedonia',
               'Nicaragua', 'Namibia', 'Moldova', 'Chad', 'Equatorial Guinea',
               'Congo, Rep.', 'Rwanda', 'Haiti', 'Kyrgyz Republic', 'Tajikistan',
               'Kosovo', 'Malawi', 'Mauritania', 'Monaco', 'Isle of Man',
               'Liechtenstein', 'Guam', 'Maldives', 'Fiji', 'Montenegro',
               'Cayman Islands', 'Togo', 'Barbados', 'Eswatini', 'Guyana',
               'Suriname', 'Sierra Leone', 'Virgin Islands (U.S.)', 'Djibouti',
               'Andorra', 'Curaçao', 'Liberia', 'Aruba', 'Greenland', 'Burundi',
               'Faroe Islands', 'Lesotho', 'Bhutan', 'Central African Republic',
               'St. Lucia', 'Cabo Verde', 'Belize', 'Gambia, The',
               'Antigua and Barbuda', 'Seychelles', 'Timor-Leste', 'San Marino',
               'Solomon Islands', 'Guinea-Bissau', 'Northern Mariana Islands',
               'Grenada', 'Comoros', 'St. Kitts and Nevis',
               'Turks and Caicos Islands', 'Vanuatu', 'Samoa',
               'St. Vincent and the Grenadines', 'American Samoa', 'Dominica',
               'Tonga', 'São Tomé and Príncipe', 'Micronesia, Fed. Sts.', 'Pala
u',
               'Marshall Islands', 'Kiribati', 'Nauru', 'Tuvalu', 'Bermuda',
               'British Virgin Islands', 'Channel Islands', 'Eritrea',
               'French Polynesia', 'Gibraltar', "Korea, Dem. People's Rep.",
               'New Caledonia', 'Sint Maarten (Dutch part)', 'South Sudan',
               'St. Martin (French part)', 'Syrian Arab Republic',
               'Venezuela, RB', 'Somalia', nan, 'World', 'East Asia & Pacific',
               'Europe & Central Asia', 'Latin America & Caribbean',
               'Middle East & North Africa', 'North America', 'South Asia',
               'Sub-Saharan Africa', 'Low income', 'Lower middle income',
               'Upper middle income', 'High income'], dtype=object)

```

We cannot call `np.sort` on this array as it contains both strings and float values (NaN is considered a float). NaN values are useless in our case anyway, so we remove them from the array.

We use list comprehension to create a list of all country names that are not NaN values. NaN is a float so we will have to cast every name to a string and compare it with the string 'nan'. Python's sorted function is then used to sort the list.

```
In [38]: countries = gdp['country'].unique()
not_nan_countries = [c for c in countries if str(c) != 'nan']
sorted(not_nan_countries)
```

```
Out[38]: ['Afghanistan',
          'Albania',
          'Algeria',
          'American Samoa',
          'Andorra',
          'Angola',
          'Antigua and Barbuda',
          'Argentina',
          'Armenia',
          'Aruba',
          'Australia',
          'Austria',
          'Azerbaijan',
          'Bahamas, The',
          'Bahrain',
          'Bangladesh',
          'Barbados',
          'Belarus',
          'Belgium',
          'Belize',
          'Benin',
          'Bermuda',
          'Bhutan',
          'Bolivia',
          'Bosnia and Herzegovina',
          'Botswana',
          'Brazil',
          'British Virgin Islands',
          'Brunei Darussalam',
          'Bulgaria',
          'Burkina Faso',
          'Burundi',
          'Cabo Verde',
          'Cambodia',
          'Cameroon',
          'Canada',
          'Cayman Islands',
          'Central African Republic',
          'Chad',
          'Channel Islands',
          'Chile',
          'China',
          'Colombia',
          'Comoros',
          'Congo, Dem. Rep.',
          'Congo, Rep.',
          'Costa Rica',
          'Croatia',
          'Cuba',
          'Curaçao',
          'Cyprus',
          'Czech Republic',
          'Côte d'Ivoire',
          'Denmark',
          'Djibouti',
          'Dominica',
          'Dominican Republic',
          'East Asia & Pacific',
          'Ecuador',
          'Egypt, Arab Rep.',
          'El Salvador',
          'Equatorial Guinea',
          'Eritrea',
          'Estonia',
```


'Eswatini',
'Ethiopia',
'Europe & Central Asia',
'Faroe Islands',
'Fiji',
'Finland',
'France',
'French Polynesia',
'Gabon',
'Gambia, The',
'Georgia',
'Germany',
'Ghana',
'Gibraltar',
'Greece',
'Greenland',
'Grenada',
'Guam',
'Guatemala',
'Guinea',
'Guinea-Bissau',
'Guyana',
'Haiti',
'High income',
'Honduras',
'Hong Kong SAR, China',
'Hungary',
'Iceland',
'India',
'Indonesia',
'Iran, Islamic Rep.',
'Iraq',
'Ireland',
'Isle of Man',
'Israel',
'Italy',
'Jamaica',
'Japan',
'Jordan',
'Kazakhstan',
'Kenya',
'Kiribati',
'Korea, Dem. People's Rep.',
'Korea, Rep.',
'Kosovo',
'Kuwait',
'Kyrgyz Republic',
'Lao PDR',
'Latin America & Caribbean',
'Latvia',
'Lebanon',
'Lesotho',
'Liberia',
'Libya',
'Liechtenstein',
'Lithuania',
'Low income',
'Lower middle income',
'Luxembourg',
'Macao SAR, China',
'Madagascar',
'Malawi',
'Malaysia',
'Maldives',

'Mali',
'Malta',
'Marshall Islands',
'Mauritania',
'Mauritius',
'Mexico',
'Micronesia, Fed. Sts.',
'Middle East & North Africa',
'Moldova',
'Monaco',
'Mongolia',
'Montenegro',
'Morocco',
'Mozambique',
'Myanmar',
'Namibia',
'Nauru',
'Nepal',
'Netherlands',
'New Caledonia',
'New Zealand',
'Nicaragua',
'Niger',
'Nigeria',
'North America',
'North Macedonia',
'Northern Mariana Islands',
'Norway',
'Oman',
'Pakistan',
'Palau',
'Panama',
'Papua New Guinea',
'Paraguay',
'Peru',
'Philippines',
'Poland',
'Portugal',
'Puerto Rico',
'Qatar',
'Romania',
'Russian Federation',
'Rwanda',
'Samoa',
'San Marino',
'Saudi Arabia',
'Senegal',
'Serbia',
'Seychelles',
'Sierra Leone',
'Singapore',
'Sint Maarten (Dutch part)',
'Slovak Republic',
'Slovenia',
'Solomon Islands',
'Somalia',
'South Africa',
'South Asia',
'South Sudan',
'Spain',
'Sri Lanka',
'St. Kitts and Nevis',
'St. Lucia',
'St. Martin (French part)',

```
'St. Vincent and the Grenadines',
'Sub-Saharan Africa',
'Sudan',
'Suriname',
'Sweden',
'Switzerland',
'Syrian Arab Republic',
'São Tomé and Príncipe',
'Tajikistan',
'Tanzania',
'Thailand',
'Timor-Leste',
'Togo',
'Tonga',
'Trinidad and Tobago',
'Tunisia',
'Turkey',
'Turkmenistan',
'Turks and Caicos Islands',
'Tuvalu',
'Uganda',
'Ukraine',
'United Arab Emirates',
'United Kingdom',
'United States',
'Upper middle income',
'Uruguay',
'Uzbekistan',
'Vanuatu',
'Venezuela, RB',
'Vietnam',
'Virgin Islands (U.S.)',
'West Bank and Gaza',
'World',
'Yemen, Rep.',
'Zambia',
'Zimbabwe']
```

Unlike before, we have fewer country names to shorten. Based on the list, the only country names different from the conventional versions that we are using for this assignment is `Brunei Darussalam`, `Timor-Leste` and `Lao PDR`.

We can then reuse the code from before (albeit with changes to `renames`) and obtain a DataFrame called `gdp_sea` that contains GDP data for Southeast Asian countries only.

```
In [39]: renames_gdp = {'Brunei Darussalam': 'Brunei',
                        'Timor-Leste': 'East Timor',
                        'Lao PDR': 'Laos'
                        }

gdp['country'] = gdp['country'].replace(renames_gdp)
gdp_sea = gdp[gdp['country'].isin(sea)].reset_index(drop = True)

gdp_sea
```

Out[39]:

	country code	Ranking	country	GDP (in millions USD)
0	IDN	16	Indonesia	1,119,191
1	THA	22	Thailand	543,650
2	PHL	33	Philippines	376,796
3	SGP	34	Singapore	372,063
4	MYS	36	Malaysia	364,702
5	VNM	45	Vietnam	261,921
6	MMR	71	Myanmar	76,086
7	KHM	103	Cambodia	27,089
8	LAO	117	Laos	18,174
9	BRN	133	Brunei	13,469
10	TLS	182	East Timor	1,674

The original DataFrame is sorted by GDP ranking instead of country name. Using `sort_values`, we sort the DataFrame by alphabetical order of the country names.

```
In [40]: gdp_sea = gdp_sea.sort_values('country')
gdp_sea
```

Out[40]:

	country code	Ranking	country	GDP (in millions USD)
9	BRN	133	Brunei	13,469
7	KHM	103	Cambodia	27,089
10	TLS	182	East Timor	1,674
0	IDN	16	Indonesia	1,119,191
8	LAO	117	Laos	18,174
4	MYS	36	Malaysia	364,702
6	MMR	71	Myanmar	76,086
2	PHL	33	Philippines	376,796
3	SGP	34	Singapore	372,063
1	THA	22	Thailand	543,650
5	VNM	45	Vietnam	261,921

The DataFrame is now in alphabetical order. We reset the indexes so they start from 0, and use `filter` to retain only the `country` and `GDP (in millions USD)` columns, as `GDP (in millions USD)` is the new Series of data that we want to merge into our final DataFrame.

```
In [41]: gdp_sea = gdp_sea.reset_index(drop = True).filter(['country', 'GDP (in mill:
gdp_sea
```

Out[41]:

	country	GDP (in millions USD)
0	Brunei	13,469
1	Cambodia	27,089
2	East Timor	1,674
3	Indonesia	1,119,191
4	Laos	18,174
5	Malaysia	364,702
6	Myanmar	76,086
7	Philippines	376,796
8	Singapore	372,063
9	Thailand	543,650
10	Vietnam	261,921

Almost there! Let us look at the data type of `GDP (in millions USD)`.

In [42]: `gdp_sea.dtypes`

Out[42]:

country	object
GDP (in millions USD)	object

dtype: object

As expected, they are strings. We will have to convert them to integers.

Just like before, the `apply` function is used in tandem with a lambda expression. The strings now contain commas instead of whitespaces to separate every three digits, so we replace the commas with empty strings.

The strings can then be casted to integers using `to_numeric`.

In [43]:

```
gdp_sea['GDP (in millions USD)'] = gdp_sea['GDP (in millions USD)'].apply(lambda x: x.replace(',', ''))
gdp_sea['GDP (in millions USD)'] = pd.to_numeric(gdp_sea['GDP (in millions USD)'], errors='coerce')
gdp_sea
```

Out[43]:

	country	GDP (in millions USD)
0	Brunei	13469
1	Cambodia	27089
2	East Timor	1674
3	Indonesia	1119191
4	Laos	18174
5	Malaysia	364702
6	Myanmar	76086
7	Philippines	376796
8	Singapore	372063
9	Thailand	543650
10	Vietnam	261921

To confirm that the casts have been performed correctly:

In [44]: `gdp_sea.dtypes`

Out[44]:

country	object
GDP (in millions USD)	int64
dtype:	object

Now we can merge the data for GDP to `life_exp_sea_pop`. We will call this DataFrame `life_exp_sea_pop_gdp`.

In [45]: `life_exp_sea_pop_gdp = life_exp_sea_pop.merge(gdp_sea)`
`life_exp_sea_pop_gdp`

Out[45]:

	country	Status	max_life_expectancy	mean_life_expectancy	mean_BMI	mean_income
0	Brunei	Developing	78.3	76.48750	33.442857	
1	Cambodia	Developing	68.7	64.34375	15.362500	
2	East Timor	Developing	68.3	64.75625	14.550000	
3	Indonesia	Developing	69.1	67.55625	21.120000	
4	Laos	Developing	65.7	62.38125	16.057143	
5	Malaysia	Developing	75.0	73.75625	32.742857	
6	Myanmar	Developing	66.6	64.20000	18.100000	
7	Philippines	Developing	68.5	67.57500	21.592857	
8	Singapore	Developed	87.0	81.47500	31.069231	
9	Thailand	Developing	74.9	73.08125	25.946154	
10	Vietnam	Developing	76.0	74.77500	13.438462	

Now that we have for the population of Southeast Asian countries as well as their GDP in 2019, we can now compute an additional metric: GDP per capita. The formula is simply GDP divided by population.

```
In [46]: life_exp_sea_pop_gdp['perCapitaGDP'] = life_exp_sea_pop_gdp['GDP (in million USD)'] / life_exp_sea_pop_gdp['population']
```

```
Out[46]:
```

	country	Status	max_life_expectancy	mean_life_expectancy	mean_BMI	mean_income
0	Brunei	Developing	78.3	76.48750	33.442857	
1	Cambodia	Developing	68.7	64.34375	15.362500	
2	East Timor	Developing	68.3	64.75625	14.550000	
3	Indonesia	Developing	69.1	67.55625	21.120000	
4	Laos	Developing	65.7	62.38125	16.057143	
5	Malaysia	Developing	75.0	73.75625	32.742857	
6	Myanmar	Developing	66.6	64.20000	18.100000	
7	Philippines	Developing	68.5	67.57500	21.592857	
8	Singapore	Developed	87.0	81.47500	31.069231	
9	Thailand	Developing	74.9	73.08125	25.946154	
10	Vietnam	Developing	76.0	74.77500	13.438462	

Our GDP columns was in millions of USD. Upon dividing by population, we get float values with too many leading zeroes. This makes the data hard to interpret. We can simply multiply each value by 1 million to get the exact figures. We will also rename the column for clarity.

```
In [47]: life_exp_sea_pop_gdp['perCapitaGDP'] = life_exp_sea_pop_gdp['perCapitaGDP'] * 1000000
life_exp_sea_pop_gdp.rename(columns = {'perCapitaGDP': 'perCapitaGDP (in USD)'})
life_exp_sea_pop_gdp
```

Out[47]:

	country	Status	max_life_expectancy	mean_life_expectancy	mean_BMI	mean_income
0	Brunei	Developing	78.3	76.48750	33.442857	
1	Cambodia	Developing	68.7	64.34375	15.362500	
2	East Timor	Developing	68.3	64.75625	14.550000	
3	Indonesia	Developing	69.1	67.55625	21.120000	
4	Laos	Developing	65.7	62.38125	16.057143	
5	Malaysia	Developing	75.0	73.75625	32.742857	
6	Myanmar	Developing	66.6	64.20000	18.100000	
7	Philippines	Developing	68.5	67.57500	21.592857	
8	Singapore	Developed	87.0	81.47500	31.069231	
9	Thailand	Developing	74.9	73.08125	25.946154	
10	Vietnam	Developing	76.0	74.77500	13.438462	

But now there are way too many decimal places! We can truncate them using `round`, a **pandas** function that rounds numbers to the desired number of decimal places. For readability, we are going with 2 decimal places and will apply it to all float values in the DataFrame.

We will iterate through every column using a for loop. If the datatype of the column is a **NumPy** float, the `round` function is applied.

After this cell, we get the final DataFrame.

```
In [48]: for col in life_exp_sea_pop_gdp:
          if life_exp_sea_pop_gdp[col].dtype == np.float64:
              life_exp_sea_pop_gdp[col] = life_exp_sea_pop_gdp[col].round(2)

life_exp_sea_pop_gdp
```


Out[48]:

	country	Status	max_life_expectancy	mean_life_expectancy	mean_BMI	mean_income
0	Brunei	Developing	78.3	76.49	33.44	
1	Cambodia	Developing	68.7	64.34	15.36	
2	East Timor	Developing	68.3	64.76	14.55	
3	Indonesia	Developing	69.1	67.56	21.12	
4	Laos	Developing	65.7	62.38	16.06	
5	Malaysia	Developing	75.0	73.76	32.74	
6	Myanmar	Developing	66.6	64.20	18.10	
7	Philippines	Developing	68.5	67.58	21.59	
8	Singapore	Developed	87.0	81.47	31.07	
9	Thailand	Developing	74.9	73.08	25.95	
10	Vietnam	Developing	76.0	74.78	13.44	

Interpreting the data

Of the 11 Southeast Asian countries in this study, Singapore is the only country with `Developed` status. Hence, it is no surprise that Singapore has the highest values of `max_life_expectancy` (87.0) and `mean_life_expectancy` (81.47). In contrast, Laos has lowest values of both `max_life_expectancy` (65.7) and `mean_life_expectancy` (62.38). Interestingly, it can be noted that countries in the centre of Southeast Asia (eg. Singapore, Brunei, Malaysia) generally have higher life expectancies than those in the north and south (eg. Laos, Indonesia, East Timor). For the north, this can probably be explained by regular instability, such as famine and war. Indonesia takes up most of the south, so its large population may contribute to the calculation of a lower life expectancy.

Based on information from [the NHS](#), the values of `mean_BMI` calculated from this dataset are quite unrealistic. It is highly likely the `BMI` values in the original dataset are inaccurate. For example, the mean BMI of the Vietnamese population is 13.44, a severely low figure.

`mean_income_composition_of_resources` seems to have strong correlation with `mean_schooling` and `perCapitaGDP`, as well as `mean_life_expectancy`. This makes sense as higher schooling increases the number of educated citizens, which in turn increases the mean income composition of resources. High income composition of resources will then lead to higher GDP per capita, giving the citizens of the country better quality of life and thus longer life expectancy.

Question 1

This question requires us to visualise expected life expectancies in Southeast Asia

based on country status (Developing or Developed). We use **Seaborn** to create a barplot for this, as **Seaborn** makes it easier for us to display the values above each bar. Furthermore, **Seaborn** automatically takes the mean of all mean life expectancies of developing countries and uses it as the value for the bar chart. Thus, no extra calculation is needed.

Since **Seaborn** is built on top of **Matplotlib**, we can use **Matplotlib**'s `figure` function to create a blank figure with a specified size, the `title` function to set the title of the graph, then call **Seaborn**'s `barplot` function to plot the bars. `ci`, or confidence intervals, is set to `None` as we do not need to display them in this scenario.

When the bar chart is plotted, the values are stored in an object created by **Matplotlib** called `containers`. We can then display the values of each bar by passing the first element of `containers` into the `bar_label` function.

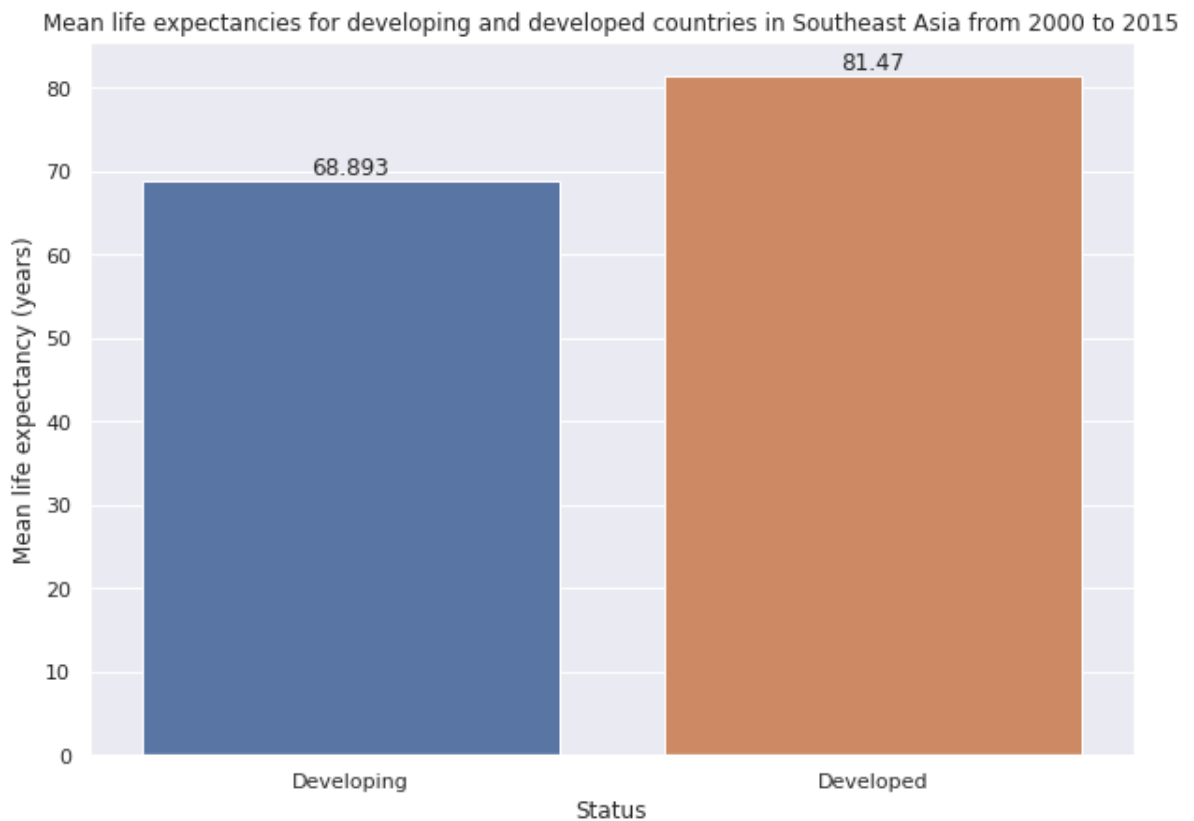
Finally, `set_ylabel` is used to display a clear y-axis label with the unit in parantheses.

```
In [49]: # this code sets the visual theme for all plots to Seaborn's default theme :
# Seaborn's theme has dark grids and softer colours which makes the plot mo
sns.set_theme()
```

```
In [50]: q1 = plt.figure(figsize = (10, 7))
q1 = plt.title('Mean life expectancies for developing and developed countrie
q1 = sns.barplot(x = 'Status', y = 'mean_life_expectancy', data = life_exp_s
q1.bar_label(q1.containers[0])
q1.set_ylabel('Mean life expectancy (years)')

q1
```

```
Out[50]: <AxesSubplot:title={'center':'Mean life expectancies for developing and d
eveloped countries in Southeast Asia from 2000 to 2015'}, xlabel='Statu
s', ylabel='Mean life expectancy (years)'
```



Insight

A bar chart is chosen to visualise and compare the expected life expectancies for developing and developed countries in Southeast Asia, since the data involved is categorical. A bar chart can give us a straightforward comparison of the data. The mean was chosen as the aggregation to represent expected life expectancies as it is a measure of central tendency. Values are more likely to occur the closer they are to the mean, so it gives us an idea of what values to 'expect' in a sample or population.

The 'mean life expectancy for developing countries' is the mean of all mean life expectancies for developing countries as shown in the DataFrame

`life_exp_sea_pop_gdp`. Since Singapore is the only developed country in Southeast Asia, the 'mean life expectancy for developed countries' is simply the mean life expectancy for Singapore. From the graph, the former value is 68.893 years whereas the latter is 81.47 years.

Thus, the insight we gain is that the difference between life expectancies for developing and developed countries in Southeast Asia is quite large (about 12.58 years). However, due to the fact that there are 10 developing countries and only one developed country, it is not a fair representation of the difference in life expectancies for developing and developed countries in general.

Question 2

Now we need to create a bar graph with countries on the x-axis, and side-by-side bars for population, mean life expectancy and adult mortality. We have dropped the data for adult mortality earlier on, but we can restore it easily. If you remember, we have

kept a copy of the original DataFrame in `life_exp_ori`. Using the `concat` function, we can slice the `country` and `Adult Mortality` columns from `life_exp_ori` and concatenate them to get a 2-column DataFrame.

The approach to handling this DataFrame is similar to what we did before. We use `replace` with the `renames` dictionary to shorten the country names, and remove all rows except those for Southeast Asian countries. We then audit the data using `info` and `describe`.

```
In [51]: adult_mort = pd.concat([life_exp_ori['country'], life_exp_ori['Adult Mortality']], axis=1)
adult_mort['country'] = adult_mort['country'].replace(renames)
adult_mort = adult_mort[adult_mort['country'].isin(sea)]

adult_mort.info()
adult_mort.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 176 entries, 383 to 2874
Data columns (total 2 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   country         176 non-null   object  
 1   Adult Mortality 176 non-null   float64
dtypes: float64(1), object(1)
memory usage: 4.1+ KB
```

Out[51]: **Adult Mortality**

count	176.000000
mean	148.846591
std	75.842350
min	1.000000
25%	84.000000
50%	158.000000
75%	213.250000
max	296.000000

There are no missing values but there is definitely some erroneous data, as evident from the minimum value of 1.

We can use the same strategy we used to remove the erroneous `BMI` values from `life_exp_sea`, but adult mortality is a metric that can be very different in each country, and it can change in large intervals. It is not easy to find the erroneous data.

A quick and dirty way is to replace all outliers with the mean of non-outlier values for each country. Outliers can be defined as any value below the lower quartile subtracted by 1.5 times the interquartile range, or above the upper quartile subtracted by 1.5 times the interquartile range*.

Since `Adult Mortality` will have to be used **Question 3**, we save a

* Source: <https://www.thoughtco.com/what-is-an-outlier-3126227>

```
In [52]: for c in adult_mort['country'].unique():
    upper = adult_mort[adult_mort['country'] == c].quantile(0.75)['Adult Mortality']
    lower = adult_mort[adult_mort['country'] == c].quantile(0.25)['Adult Mortality']
    outlier = 1.5 * (upper - lower)

    adult_mort.loc[(adult_mort['country'] == c) & (adult_mort['Adult Mortality'] > upper), 'Adult Mortality'] = upper
    adult_mort.loc[(adult_mort['country'] == c) & (adult_mort['Adult Mortality'] < lower), 'Adult Mortality'] = lower

    mean_am = adult_mort[adult_mort['country'] == c]['Adult Mortality'].mean()
    adult_mort.loc[(adult_mort['country'] == c) & (adult_mort['Adult Mortality'] > mean_am), 'Adult Mortality'] = mean_am

adult_mort
```

Out[52]:

	country	Adult Mortality
383	Brunei	85.750000
382	Brunei	85.750000
381	Brunei	95.000000
380	Brunei	89.000000
379	Brunei	89.000000
...
2878	Vietnam	131.000000
2877	Vietnam	134.133333
2876	Vietnam	129.000000
2875	Vietnam	128.000000
2874	Vietnam	127.000000

176 rows × 2 columns

We then create a `groupby` object based on `country` and aggregate the values of `Adult Mortality` to mean values. Finally, `Adult Mortality` is renamed to `mean_adult_mortality`.

```
In [53]: adult_mort_agg = adult_mort.groupby(['country'], as_index = False).agg('mean')
    adult_mort_agg.rename(columns = {'Adult Mortality': 'mean_adult_mortality'})
```

Now that `life_exp_sea_pop_gdp` and `adult_mort_agg` share an identical `country` column, they can be merged easily. We assign this new DataFrame to the variable `q2_df` so the unmerged version is not overwritten.

```
In [54]: q2_df = life_exp_sea_pop_gdp.merge(adult_mort_agg)

q2_df
```

Out [54]:

	country	Status	max_life_expectancy	mean_life_expectancy	mean_BMI	mean_income
0	Brunei	Developing	78.3	76.49	33.44	
1	Cambodia	Developing	68.7	64.34	15.36	
2	East Timor	Developing	68.3	64.76	14.55	
3	Indonesia	Developing	69.1	67.56	21.12	
4	Laos	Developing	65.7	62.38	16.06	
5	Malaysia	Developing	75.0	73.76	32.74	
6	Myanmar	Developing	66.6	64.20	18.10	
7	Philippines	Developing	68.5	67.58	21.59	
8	Singapore	Developed	87.0	81.47	31.07	
9	Thailand	Developing	74.9	73.08	25.95	
10	Vietnam	Developing	76.0	74.78	13.44	

In order to visualise the values for `population`, `mean_life_expectancy` and `mean_adult_mortality` in a side-by-side bar chart, we will need to unpivot these columns (ie. convert them from 'wide' to 'long'). This can be done with **pandas's** `melt` function.

Since `country` is our categorical values, it is set as the sole identifier variable, whereas the three variables we want to represent are the value variables. The 'melted' DataFrame is assigned to a new variable called `melted`.

We also rename the column labels and the names of each statistic for consistency in capitalisation and to show the units in parentheses.

```
In [55]: melted = pd.melt(q2_df, id_vars = 'country',
                        value_vars = ['population', 'mean_life_expectancy', 'mean_adult_mortality'],
                        var_name = 'Statistic',
                        value_name = 'Value',
                        inplace = True)

melted.rename(columns = {'country': 'Country',
                        'variable': 'Statistic',
                        'value': 'Value'},
              inplace = True)

melted['Statistic'].replace({'population': 'Population',
                        'mean_life_expectancy': 'Mean life expectancy (years)',
                        'mean_adult_mortality': 'Mean adult mortality rate (per 1000 live births)'},
                          inplace = True)

melted
```

Out[55]:

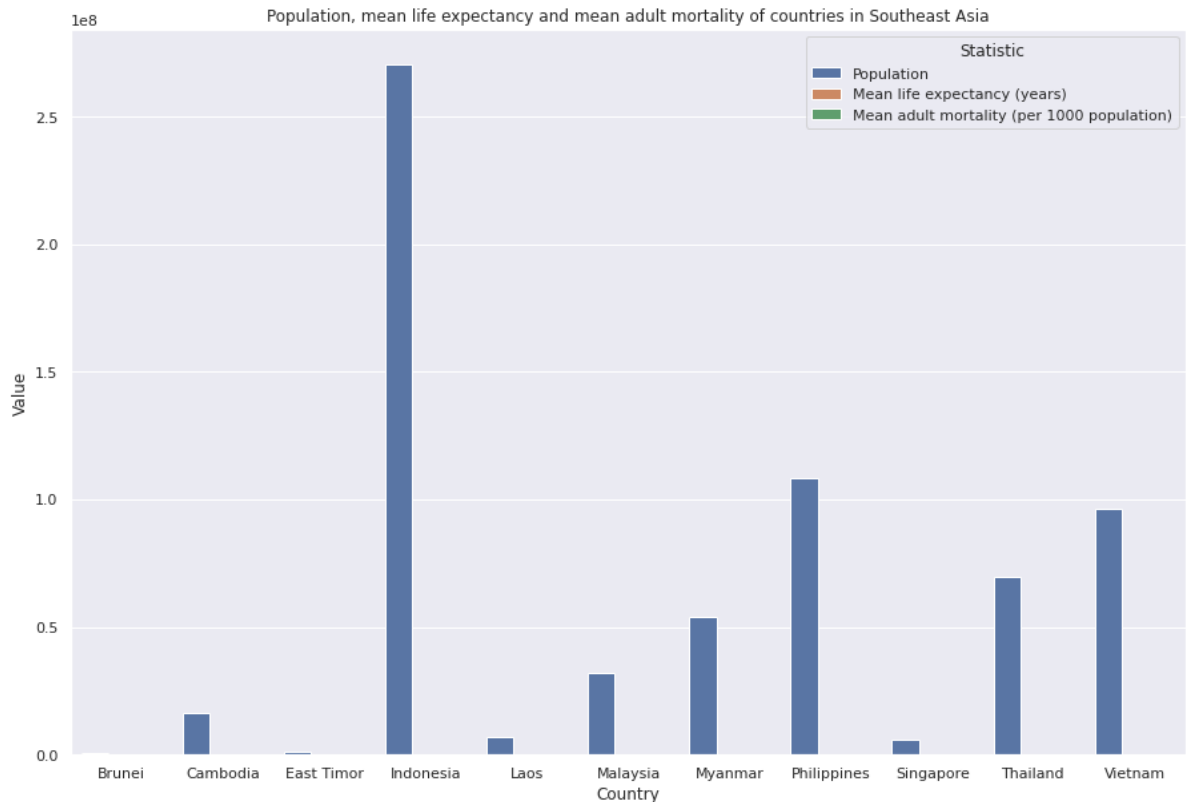
	Country	Statistic	Value
0	Brunei	Population	4.330000e+05
1	Cambodia	Population	1.648700e+07
2	East Timor	Population	1.293000e+06
3	Indonesia	Population	2.706260e+08
4	Laos	Population	7.169000e+06
5	Malaysia	Population	3.195000e+07
6	Myanmar	Population	5.404500e+07
7	Philippines	Population	1.081170e+08
8	Singapore	Population	5.804000e+06
9	Thailand	Population	6.962600e+07
10	Vietnam	Population	9.646200e+07
11	Brunei	Mean life expectancy (years)	7.649000e+01
12	Cambodia	Mean life expectancy (years)	6.434000e+01
13	East Timor	Mean life expectancy (years)	6.476000e+01
14	Indonesia	Mean life expectancy (years)	6.756000e+01
15	Laos	Mean life expectancy (years)	6.238000e+01
16	Malaysia	Mean life expectancy (years)	7.376000e+01
17	Myanmar	Mean life expectancy (years)	6.420000e+01
18	Philippines	Mean life expectancy (years)	6.758000e+01
19	Singapore	Mean life expectancy (years)	8.147000e+01
20	Thailand	Mean life expectancy (years)	7.308000e+01
21	Vietnam	Mean life expectancy (years)	7.478000e+01
22	Brunei	Mean adult mortality (per 1000 population)	8.575000e+01
23	Cambodia	Mean adult mortality (per 1000 population)	2.213571e+02
24	East Timor	Mean adult mortality (per 1000 population)	1.699091e+02
25	Indonesia	Mean adult mortality (per 1000 population)	1.856923e+02
26	Laos	Mean adult mortality (per 1000 population)	2.369231e+02
27	Malaysia	Mean adult mortality (per 1000 population)	1.323846e+02
28	Myanmar	Mean adult mortality (per 1000 population)	1.543125e+02
29	Philippines	Mean adult mortality (per 1000 population)	2.180714e+02
30	Singapore	Mean adult mortality (per 1000 population)	6.573333e+01
31	Thailand	Mean adult mortality (per 1000 population)	1.700000e+02
32	Vietnam	Mean adult mortality (per 1000 population)	1.341333e+02

Now we can use **Seaborn** to create a barplot of `melted`, where `country` is on the x-

axis, `value` is on the y-axis, and `hue` is set to `variable` so we get side-by-side multicoloured bars, each representing a different variable for that country.

```
In [56]: q2 = plt.figure(figsize = (15, 10))
q2 = plt.title('Population, mean life expectancy and mean adult mortality of countries in Southeast Asia')
q2 = sns.barplot(x = 'Country', y = 'Value', hue = 'Statistic', data = melted)
q2
```

```
Out[56]: <AxesSubplot:title={'center':'Population, mean life expectancy and mean adult mortality of countries in Southeast Asia'}, xlabel='Country', ylabel='Value'>
```



Notice how the orange and green bars are not visible. This is because the values for `population` are too large relative to that of the other two variables. To shorten this disparity between numbers, we can represent the values of `population` in millions instead.

A new column called `population (in millions)` is created, containing the values of `population` divided by 1 million. Then, we do the same melting process with `population (in millions)` replacing `population`.

```
In [57]: melted.loc[melted['Statistic'] == 'Population', 'Value'] /= 1000000
melted['Statistic'].replace('Population', 'Population (in millions)', inplace=True)
melted
```


Out [57] :

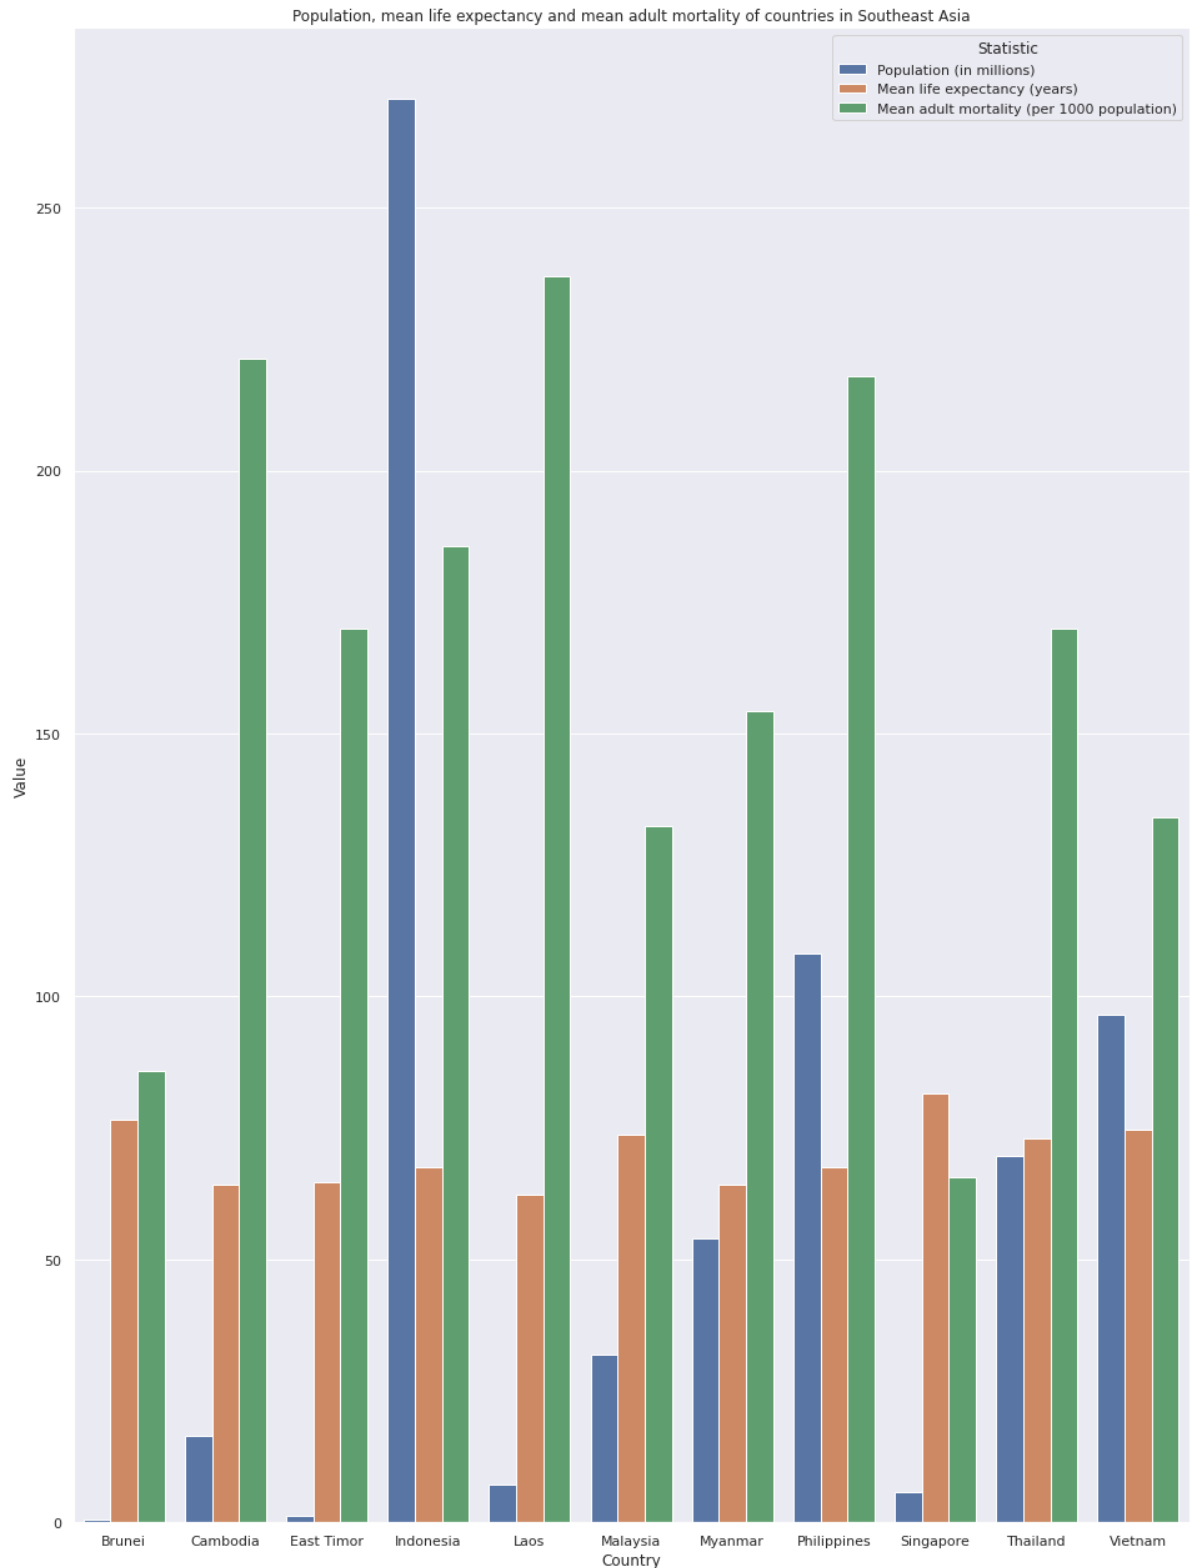
	Country	Statistic	Value
0	Brunei	Population (in millions)	0.433000
1	Cambodia	Population (in millions)	16.487000
2	East Timor	Population (in millions)	1.293000
3	Indonesia	Population (in millions)	270.626000
4	Laos	Population (in millions)	7.169000
5	Malaysia	Population (in millions)	31.950000
6	Myanmar	Population (in millions)	54.045000
7	Philippines	Population (in millions)	108.117000
8	Singapore	Population (in millions)	5.804000
9	Thailand	Population (in millions)	69.626000
10	Vietnam	Population (in millions)	96.462000
11	Brunei	Mean life expectancy (years)	76.490000
12	Cambodia	Mean life expectancy (years)	64.340000
13	East Timor	Mean life expectancy (years)	64.760000
14	Indonesia	Mean life expectancy (years)	67.560000
15	Laos	Mean life expectancy (years)	62.380000
16	Malaysia	Mean life expectancy (years)	73.760000
17	Myanmar	Mean life expectancy (years)	64.200000
18	Philippines	Mean life expectancy (years)	67.580000
19	Singapore	Mean life expectancy (years)	81.470000
20	Thailand	Mean life expectancy (years)	73.080000
21	Vietnam	Mean life expectancy (years)	74.780000
22	Brunei	Mean adult mortality (per 1000 population)	85.750000
23	Cambodia	Mean adult mortality (per 1000 population)	221.357143
24	East Timor	Mean adult mortality (per 1000 population)	169.909091
25	Indonesia	Mean adult mortality (per 1000 population)	185.692308
26	Laos	Mean adult mortality (per 1000 population)	236.923077
27	Malaysia	Mean adult mortality (per 1000 population)	132.384615
28	Myanmar	Mean adult mortality (per 1000 population)	154.312500
29	Philippines	Mean adult mortality (per 1000 population)	218.071429
30	Singapore	Mean adult mortality (per 1000 population)	65.733333
31	Thailand	Mean adult mortality (per 1000 population)	170.000000
32	Vietnam	Mean adult mortality (per 1000 population)	134.133333

The same code is used to recreate the bar graph, albeit with the figure size increased.

This will be explained later.

```
In [58]: q2 = plt.figure(figsize = (16, 22))
q2 = plt.title('Population, mean life expectancy and mean adult mortality of countries in Southeast Asia')
q2 = sns.barplot(x = 'Country', y = 'Value', hue = 'Statistic', data = melted)
q2
```

```
Out[58]: <AxesSubplot:title={'center': 'Population, mean life expectancy and mean adult mortality of countries in Southeast Asia'}, xlabel='Country', ylabel='Value'>
```



Now we have a much better bar graph that visualises the values of population, mean life expectancy and mean adult mortality.

The problem

Unfortunately, the data presented in this graph is not very appropriate. It does not really give much insight at all, and could be misleading to readers.

This is due to the difference in units. Population is a discrete and unitless statistic, whereas mean life expectancy is continuous and measured in years. Mean adult mortality is also discrete and unitless. It may seem like it can be linked closely to population, but mean adult mortality in this graph is measured per 1000 population. This is why the graph can be misleading. In order to accommodate the huge range of population values, it has to be visualised in millions while mean adult mortality remains visualised in thousands. It is tough to study the relationship between population and mean adult mortality from the graph this way.

Hence, it is not a good idea to represent these three metrics in a side-by-side bar chart. The ranges of each metric vary greatly, so adjustment of units is required. Sometimes, even adjustment of units cannot enable extreme values to be shown clearly in a standard sized figure. For example, the 22 in `figsize = (16, 22)` (ie. the height of the plot in inches) is the bare minimum required to show a miniscule bar for the population of Brunei.

An appropriate side-by-side bar chart for this dataset would be one comparing the values `mean_life_expectancy` and `max_life_expectancy`, or one that compares the values of `GDP` and `perCapitaGDP`.

Question 3

For this question, we require the non-aggregated data from `LifeExpectancyData-v2.csv`, as we will be visualising the changes in statistics over a duration of years specific to a single country - Singapore.

The non-aggregated data we need is readily available in the DataFrame `life_exp_sea`, but it is missing the `Year` column. We can use `merge` to combine the Singapore section of `life_exp_sea` with the `Year` column from `life_exp_ori`. The `Year` column alone does not share any columns with `life_exp_sea`, but both still share the same indexes since `life_exp_sea` was extracted from `life_exp`. Thus, we set both `left_index` and `right_index` as `True` so **pandas** knows to join them by using the indexes as the key. We assign this new DataFrame to a variable called `sg`.

We will be investigating `Adult Mortality` and `infant deaths`, so `sg` is merged with the `Adult Mortality` column from `adult_mort`, also using the indexes as the key.

The data for infant deaths is available in the `life_exp_ori` DataFrame. We can subset the Singapore part of this DataFrame and assign it to a temporary variable called `sg_ori` so the next line of code is not too long. `sg` is then merged with the `infant deaths` column of `sg_ori`, also on the indexes.

The rename function is then used to create some consistency with capitalisation in the column labels, and the `sg` DataFrame is complete.

```
In [59]: sg = life_exp_sea[life_exp_sea['country'] == 'Singapore'].merge(life_exp_ori,
                                                                    left_index = True,
                                                                    right_index = True)

sg = sg.merge(adult_mort['Adult Mortality'], left_index = True, right_index = True)

sg_ori = life_exp_ori[life_exp_ori['country'] == 'Singapore']
sg = sg.merge(sg_ori['infant deaths'], left_index = True, right_index = True)

sg.rename(columns = {'Adult Mortality': 'Adult mortality',
                    'infant deaths': 'Infant deaths'},
          inplace = True)

sg
```

```
Out[59]:
```

	country	Status	Life expectancy	BMI	Income composition of resources	Schooling	Year	Adult mortality
2328	Singapore	Developed	78.3	28.500000	0.810	12.5	2000	78.000000
2327	Singapore	Developed	78.7	28.900000	0.820	12.7	2001	76.000000
2326	Singapore	Developed	79.0	29.200000	0.818	12.6	2002	74.000000
2325	Singapore	Developed	79.3	29.600000	0.819	12.7	2003	73.000000
2324	Singapore	Developed	79.7	29.900000	0.820	12.7	2004	71.000000
2323	Singapore	Developed	82.0	31.069231	0.821	12.6	2005	69.000000
2322	Singapore	Developed	87.0	31.069231	0.839	13.9	2006	66.000000
2321	Singapore	Developed	81.1	31.069231	0.873	14.1	2007	65.000000
2320	Singapore	Developed	81.4	31.200000	0.880	14.2	2008	64.000000
2319	Singapore	Developed	81.7	31.500000	0.887	14.4	2009	62.000000
2318	Singapore	Developed	82.0	31.800000	0.889	14.5	2010	61.000000
2317	Singapore	Developed	82.2	32.100000	0.911	15.2	2011	65.733333
2316	Singapore	Developed	82.5	32.400000	0.917	15.4	2012	59.000000
2315	Singapore	Developed	82.7	32.700000	0.920	15.4	2013	57.000000
2314	Singapore	Developed	82.9	32.900000	0.922	15.4	2014	56.000000
2313	Singapore	Developed	83.1	33.200000	0.924	15.4	2015	55.000000

To plot a line graph of `Life expectancy` over time, we once again use **Matplotlib** to initialise a figure with a set size and title, then use **Seaborn's** `lineplot` function with `Year` on the x-axis and `Life expectancy` on the y-axis.

The values on the x-axis will be shown in intervals of two. To force the graph to show every year on the x-axis, we use `set_xticks`, which can take an array of appropriate x-axis values and display them explicitly. The argument we pass in is

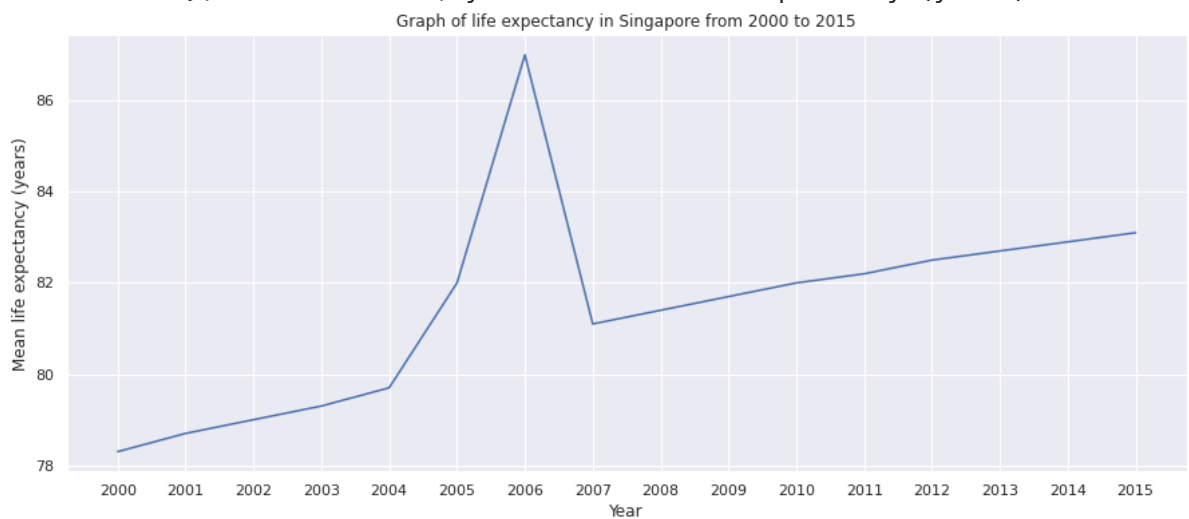
`sg['Year'].unique()` , since this returns an array of all unique values of `Year` , which is essentially every year.

Finally, the label of the y-axis is edited to include the unit in parentheses, and the graph is done.

```
In [60]: q3_1 = plt.figure(figsize = (15, 6))
q3_1 = plt.title('Graph of life expectancy in Singapore from 2000 to 2015')
q3_1 = sns.lineplot(x = 'Year', y = 'Life expectancy', data = sg)
q3_1.set_xticks(sg['Year'].unique())
q3_1.set_ylabel('Mean life expectancy (years)')

q3_1
```

```
Out[60]: <AxesSubplot:title={'center':'Graph of life expectancy in Singapore from 2000 to 2015'}, xlabel='Year', ylabel='Mean life expectancy (years)'>
```



To display two lines in a single plot, we have to make use of the `melt` function again. `melt` is applied to `sg` with `Year` as the identifier variable and `Adult mortality` and `Infant deaths` as value variables. Like in **Question 2**, the column labels and variable names are renamed for consistency and clarity.

```
In [61]: sg_melted = pd.melt(sg, id_vars = 'Year', value_vars = ['Adult mortality',
    sg_melted.rename(columns = {'country': 'Country',
    'variable': 'Statistic',
    'value': 'Value'},
    inplace = True)

    sg_melted['Statistic'].replace({'Adult mortality': 'Adult mortality (per 100
    'Infant deaths': 'Infant deaths (per 1000 po
    inplace = True)

    sg_melted
```

Out[61]:

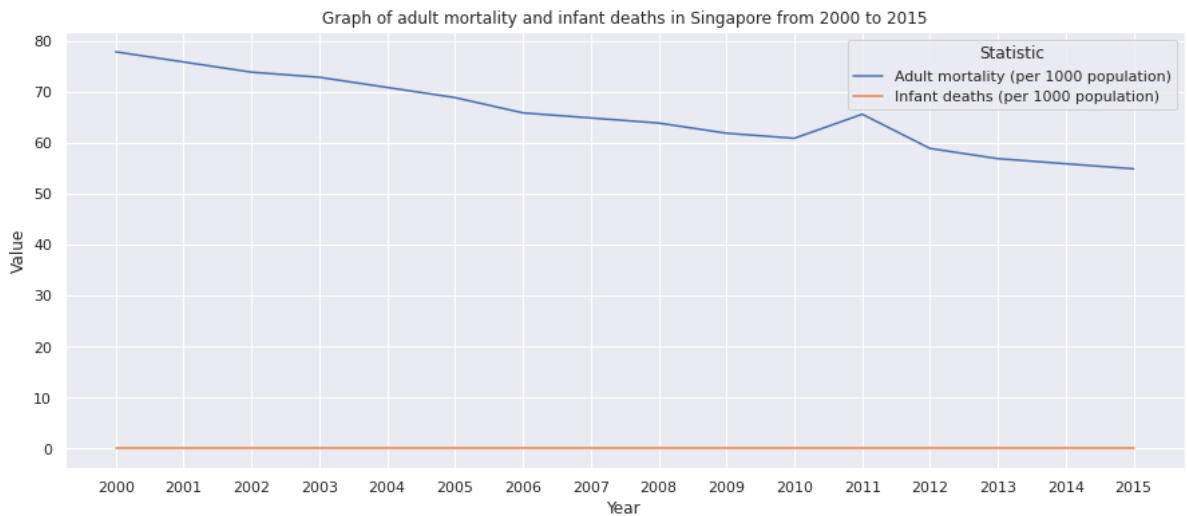
	Year	Statistic	Value
0	2000	Adult mortality (per 1000 population)	78.000000
1	2001	Adult mortality (per 1000 population)	76.000000
2	2002	Adult mortality (per 1000 population)	74.000000
3	2003	Adult mortality (per 1000 population)	73.000000
4	2004	Adult mortality (per 1000 population)	71.000000
5	2005	Adult mortality (per 1000 population)	69.000000
6	2006	Adult mortality (per 1000 population)	66.000000
7	2007	Adult mortality (per 1000 population)	65.000000
8	2008	Adult mortality (per 1000 population)	64.000000
9	2009	Adult mortality (per 1000 population)	62.000000
10	2010	Adult mortality (per 1000 population)	61.000000
11	2011	Adult mortality (per 1000 population)	65.733333
12	2012	Adult mortality (per 1000 population)	59.000000
13	2013	Adult mortality (per 1000 population)	57.000000
14	2014	Adult mortality (per 1000 population)	56.000000
15	2015	Adult mortality (per 1000 population)	55.000000
16	2000	Infant deaths (per 1000 population)	0.000000
17	2001	Infant deaths (per 1000 population)	0.000000
18	2002	Infant deaths (per 1000 population)	0.000000
19	2003	Infant deaths (per 1000 population)	0.000000
20	2004	Infant deaths (per 1000 population)	0.000000
21	2005	Infant deaths (per 1000 population)	0.000000
22	2006	Infant deaths (per 1000 population)	0.000000
23	2007	Infant deaths (per 1000 population)	0.000000
24	2008	Infant deaths (per 1000 population)	0.000000
25	2009	Infant deaths (per 1000 population)	0.000000
26	2010	Infant deaths (per 1000 population)	0.000000
27	2011	Infant deaths (per 1000 population)	0.000000
28	2012	Infant deaths (per 1000 population)	0.000000
29	2013	Infant deaths (per 1000 population)	0.000000
30	2014	Infant deaths (per 1000 population)	0.000000
31	2015	Infant deaths (per 1000 population)	0.000000

Using the melted DataFrame, can now use `lineplot` to create a graph with two lines. Like in the previous graph, the x-tick labels are manually set to show every year.

```
In [62]: q3_2 = plt.figure(figsize = (15, 6))
q3_2 = plt.title('Graph of adult mortality and infant deaths in Singapore from 2000 to 2015')
q3_2 = sns.lineplot(x = 'Year', y = 'Value', hue = 'Statistic', data = sg_merged)
q3_2.set_xticks(sg_merged['Year'].unique())

q3_2
```

```
Out[62]: <AxesSubplot:title={'center':'Graph of adult mortality and infant deaths in Singapore from 2000 to 2015'}, xlabel='Year', ylabel='Value'>
```



Explanation

The line graph titled **Graph of life expectancy in Singapore from 2000 to 2015** would be useful to study the advancement of healthcare and the welfare of citizens in Singapore from the year 2000 to 2015. These are major contributing factors of life expectancy. However, this graph alone can only provide a very general visualisation, and may not even be perfectly accurate (suggested by the outlier value for the year 2006). For a more insightful visualisation, plots for other metrics should be included as well.

Based on the graph titled **Graph of adult mortality and infant deaths in Singapore from 2000 to 2015**, life expectancy in Singapore increases as the adult mortality rate decreases. Logically, the same can be said about infant deaths, but according to this dataset, Singapore registered zero infant deaths from 2000 to 2015, meaning no relationship can technically be drawn between this and life expectancy. This can either be an impressive medical feat on Singapore's part, or the data has simply been erroneously omitted from the dataset. Other than providing insight to the growth of healthcare in Singapore, this graph can be combined with the previous graph to present a strong visualisation of how the reduction in the adult mortality rate contributes to increasing life expectancy.

Conclusion

The analysis of data from these three datasets provided us with some knowledge about the life expectancies, populations and GDP of Southeast Asian countries. As a summary, Southeast Asia is mostly still a developing region of the world, with Singapore being the sole developed country and therefore having the highest life

expectancy as well as GDP per capita. Fortunately, life expectancy in all Southeast Asian countries generally increases with time. Regardless of the rate of increase, this is a good sign that all countries in this region are seeing advancement in healthcare.

In addition, this assignment was a great opportunity to practise wrangling and visualizing real-world data for the first time in the course of this unit. It was a fun way to experience the powerful tools Python has to offer in the field of data science.